

Bio Bigdata

2014.9.30

Bio Bigdata

BICube min-kyung Kim

강남 토즈 타워점 401호

Outline

- ◆ **Bigdata**
- ◆ Hadoop
- ◆ Spark
- ◆ 3rd party Biobigdata
- ◆ Machine Learning
- ◆ Streamng
- ◆ Future



What is Bigdata?

- Data is massive, unstructured, and dirty
- Data sets so large and complex
- Difficult to process using traditional data processing applications
- Massively parallel software running on tens, hundreds, or even thousands of servers

■ What is Bigdata?

Why Biobigdata?

- Asian Genome:3.3Billion 35bp 104GB
- Afrian Genome:4.0Billion 35bp 144GB
- 1000Genomes=5,994,000 process

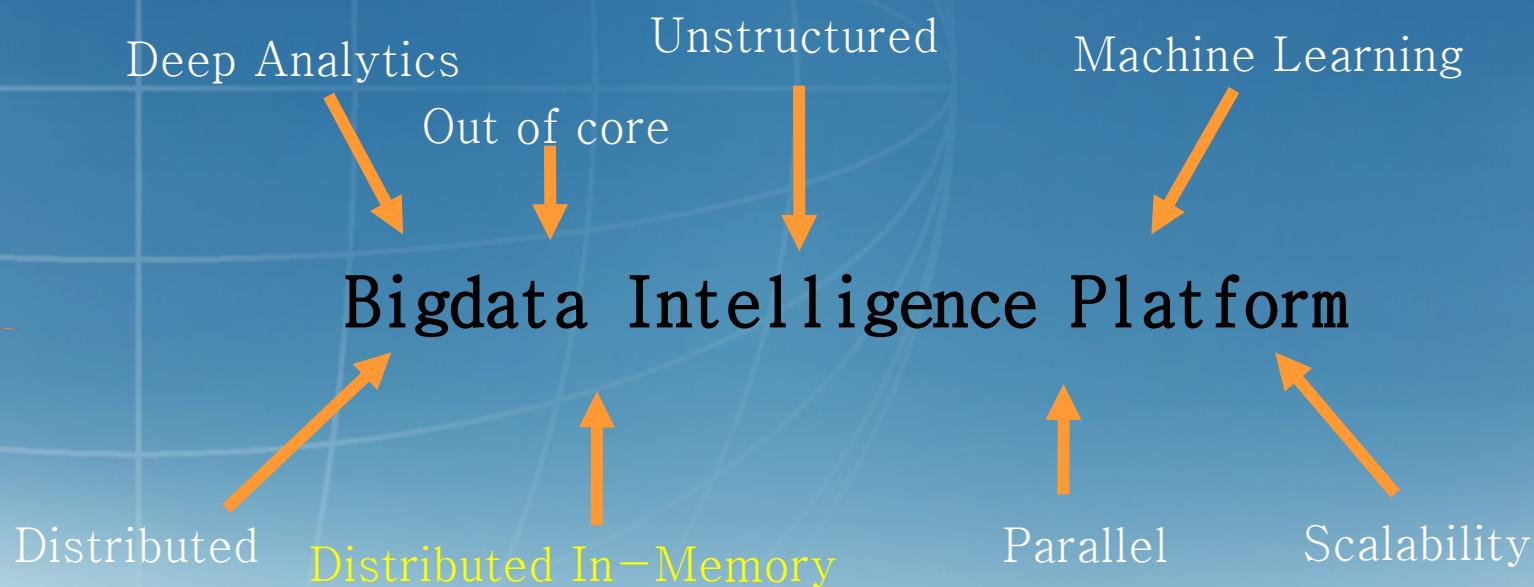
=23,976,000 hours

=2737 years

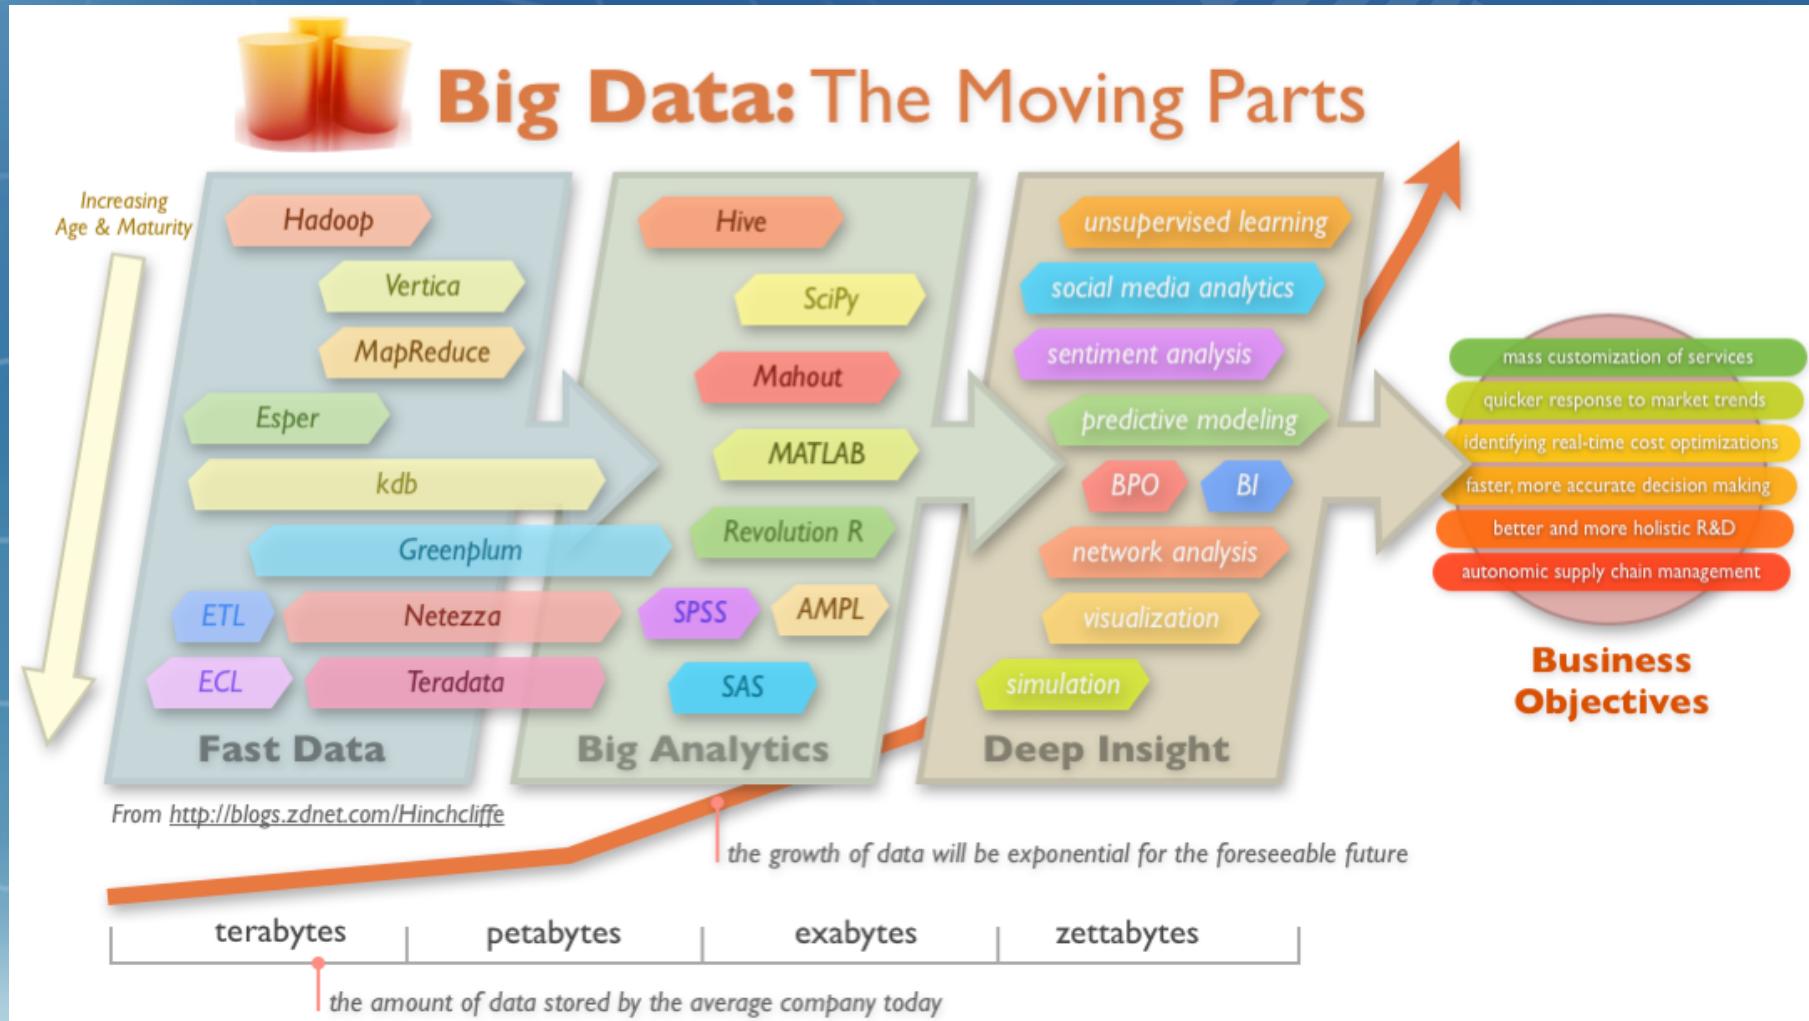
by Deepak Singh



- Rethink algorithms
- Rethink Computing
- Rethink Data Management
- Rethink Data Sharing



More Accurate Decision Making



Why Bigdata?

Central limit theorem

the arithmetic mean of a sufficiently large number of iterates of independent random variables

Overfitting, Six sigma

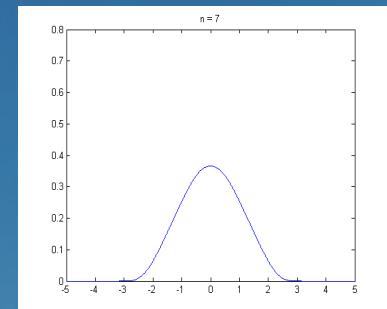
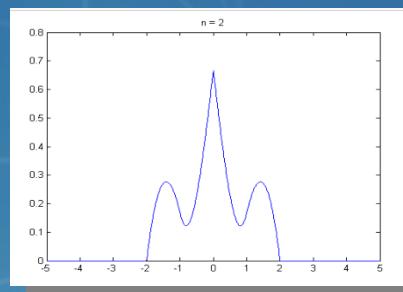
So bigdata is not necessary !

Large data is enough

But

Why Markov Assumtions?

Why Independent Assumptions?



■ Outline

- ◆ Bigdata
- ◆ **Hadoop**
- ◆ Spark
- ◆ Future



What is Hadoop?

- A open-source software for **reliable, scalable, distributed computing.**
- A framework that allows for the distributed **processing of large data sets** across clusters of computers using simple programming models
- scale up from single servers to **thousands** of machines, each offering local computation and **storage**.



Included modules(Hadoop 2)

- Hadoop Common: Utilities that support the other Hadoop modules.
- Hadoop Distributed File System (HDFS™): A distributed file system
- Hadoop YARN: A framework for job scheduling and cluster resource management.
- Hadoop MapReduce: A YARN-based system for parallel processing of large data sets.



Hadoop Distributed Filesystem(HDFS) Overviewing

- Run on Commodity Server
- The Idea of “Write once, Read many times”
- Large Streaming Reads
- High Throughput is more important than low latency



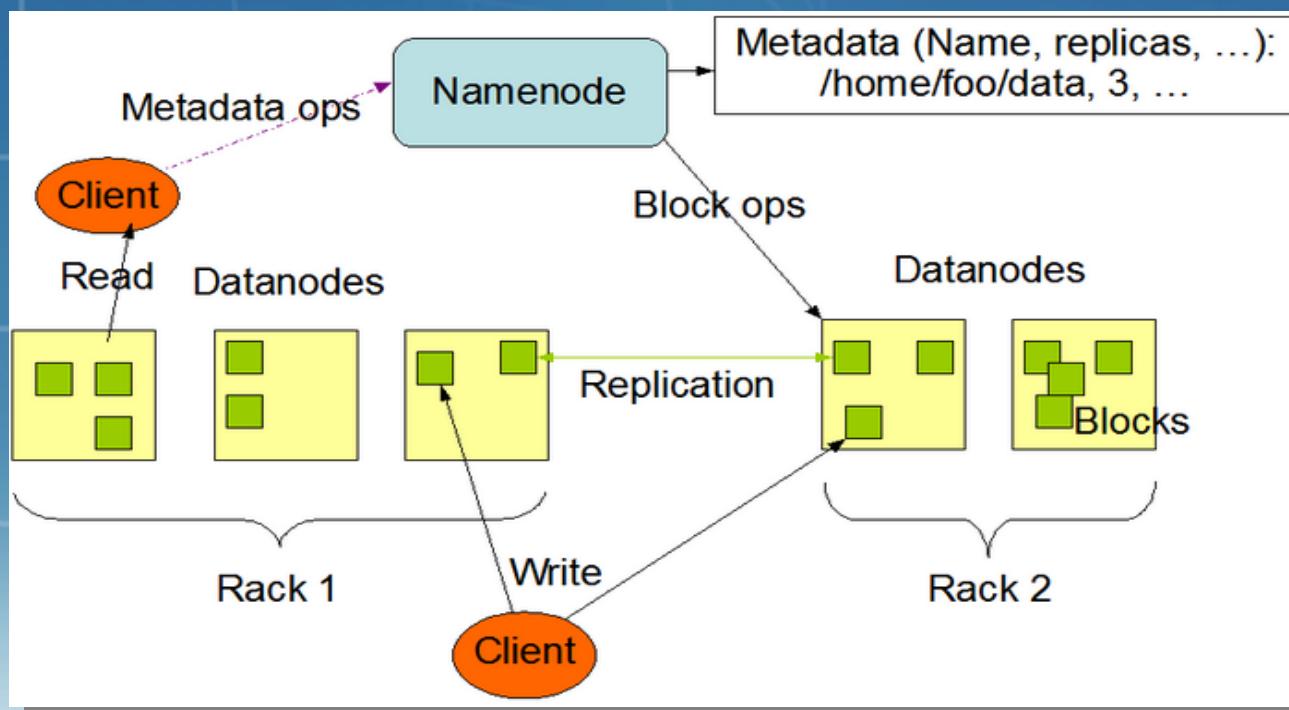
>

more important than



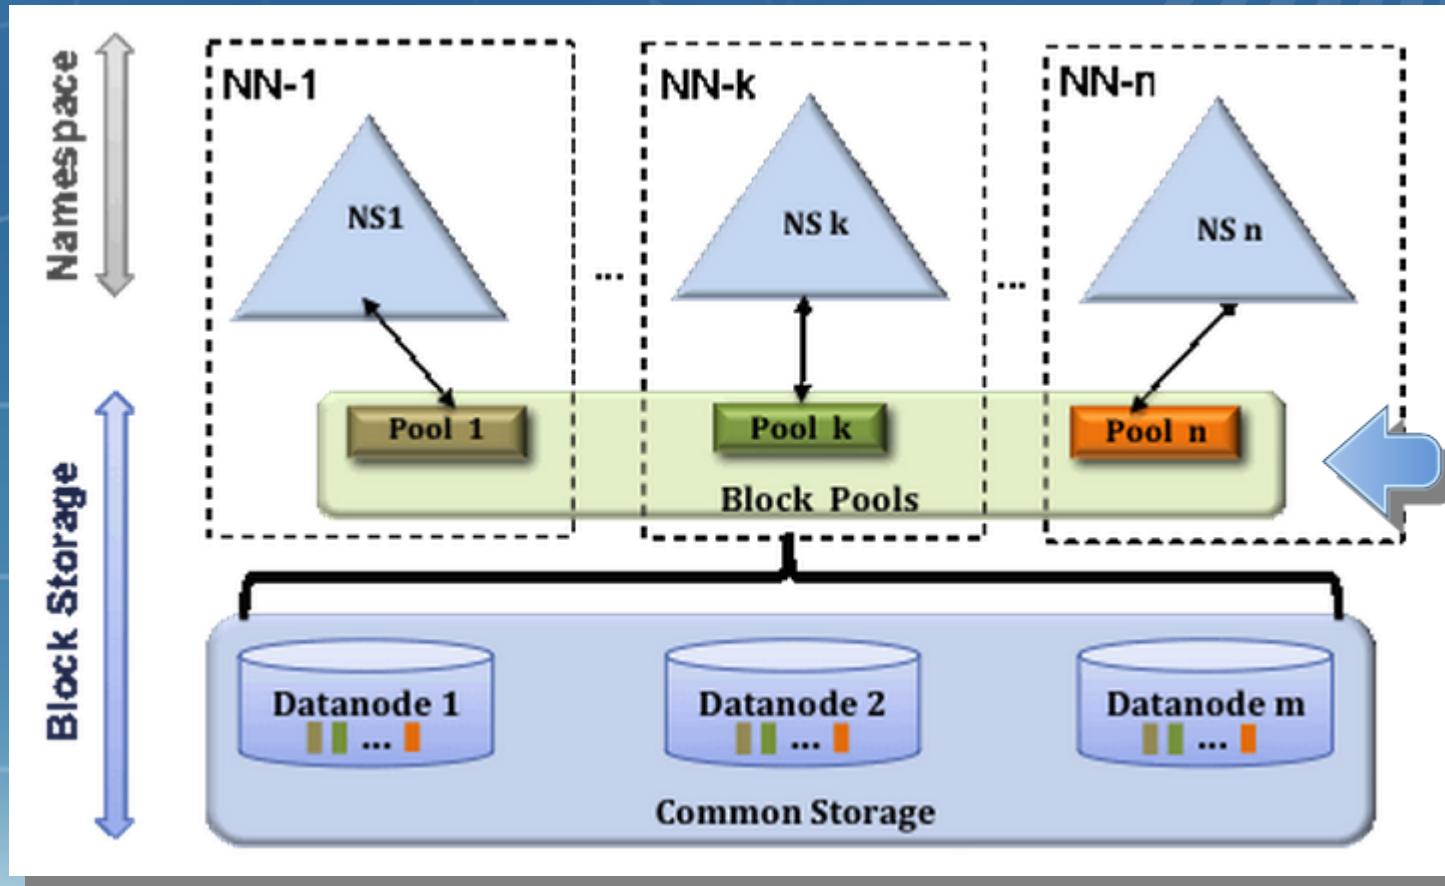
HDFS Architecture

- Stored as Block
- Provide Reliability through Replication
- NameNode stores Metadata and Manage Access
- No data caching(Large dataset)



HDFS Federation

- Block management



File Storage

- NameNode

Metadata: **filename**, Location of **block**, file **attribute**
keep metadata in RAM

Metadata Size is limited to the amount of the Name Node
RAM size

- DataNode

Store file contents as block

Same file is stored on Different DataNode

Same Block is replicated across Several DataNode

Periodically sends a report to the NameNode

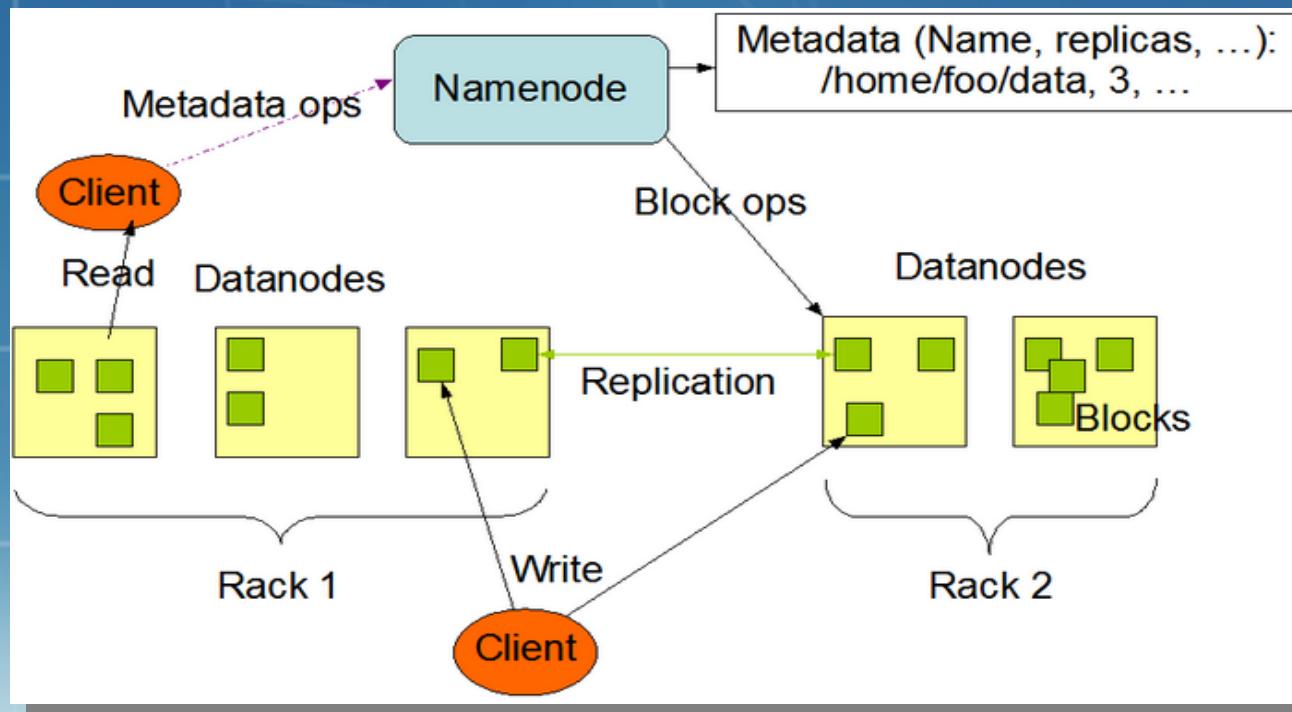
exchange heartbeats with each nodes

Client always read from nearest node



Internals of file read

- Client is guided by NameNode
- Client directly contacts DataNode
- Scale to large number of clients

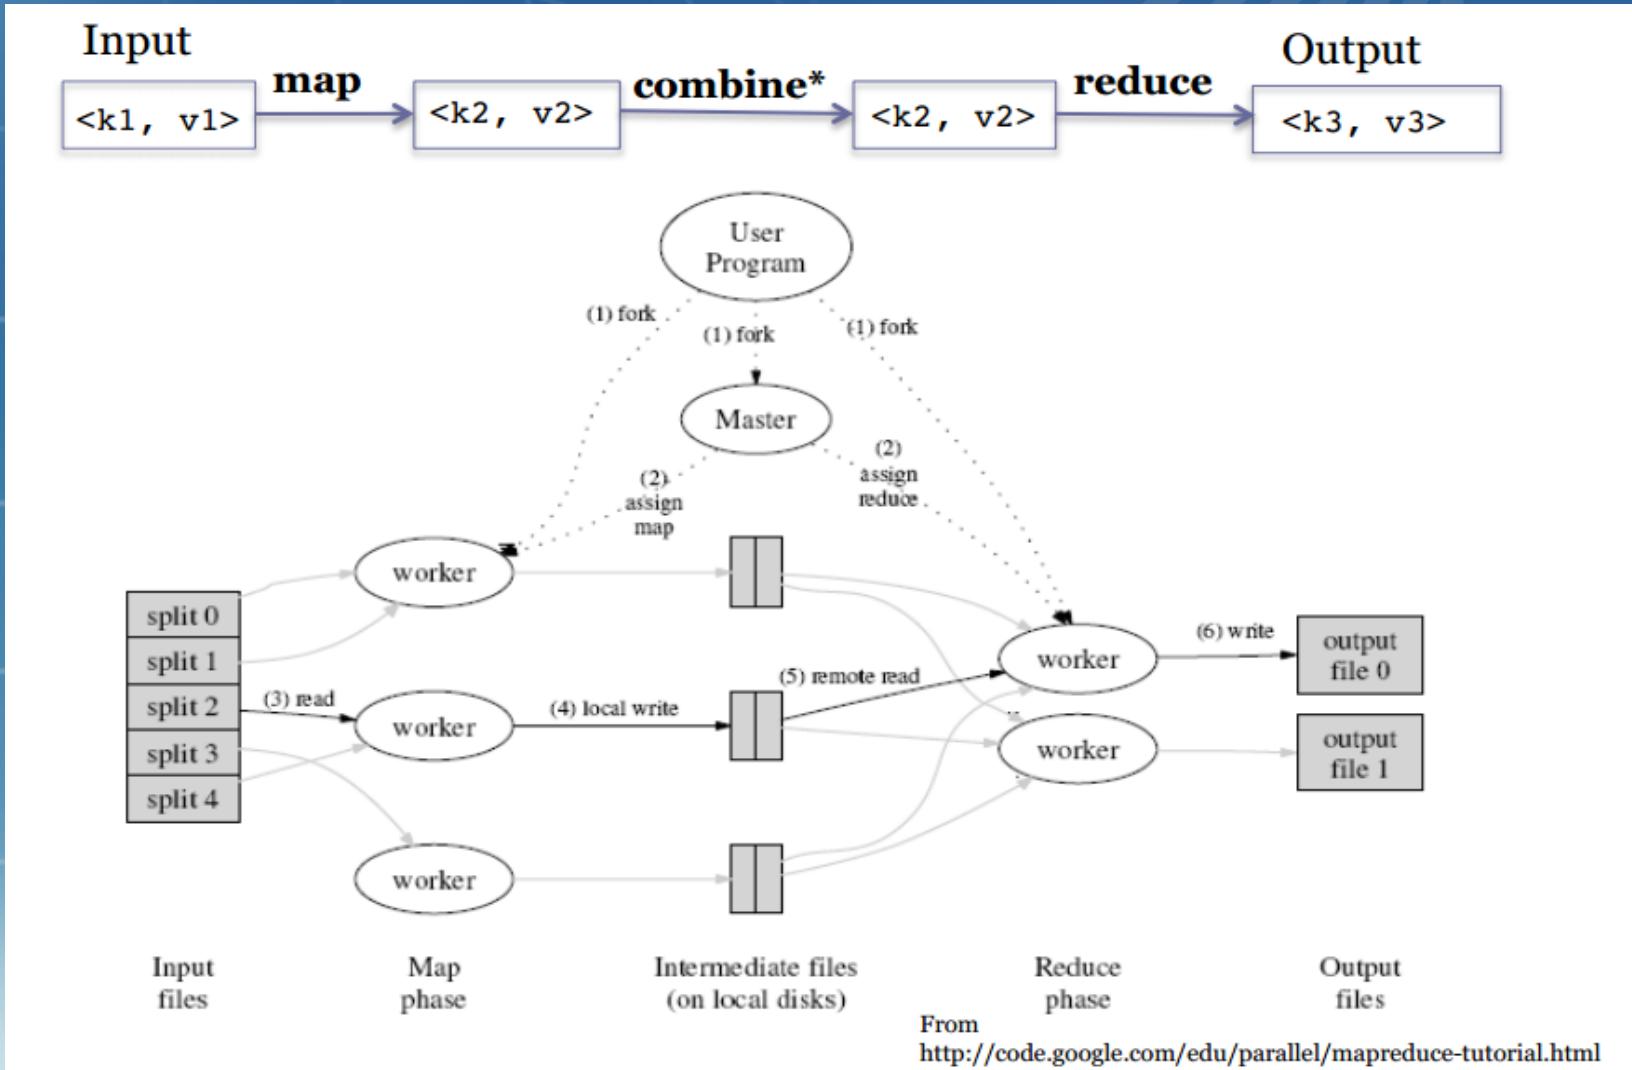


MapReduce

- Method for distributing a task
- Each node processes data on the self node
- Two phases : map, reduce
 - Map : $(K_1, V_1) \rightarrow (K_2, V_2)$
 - Reduce : $(K_2, \text{list}(V_2)) \rightarrow \text{list}(K_3, V_3)$
- Automatic parallelization
- Fault-tolerance
- Usually written in java

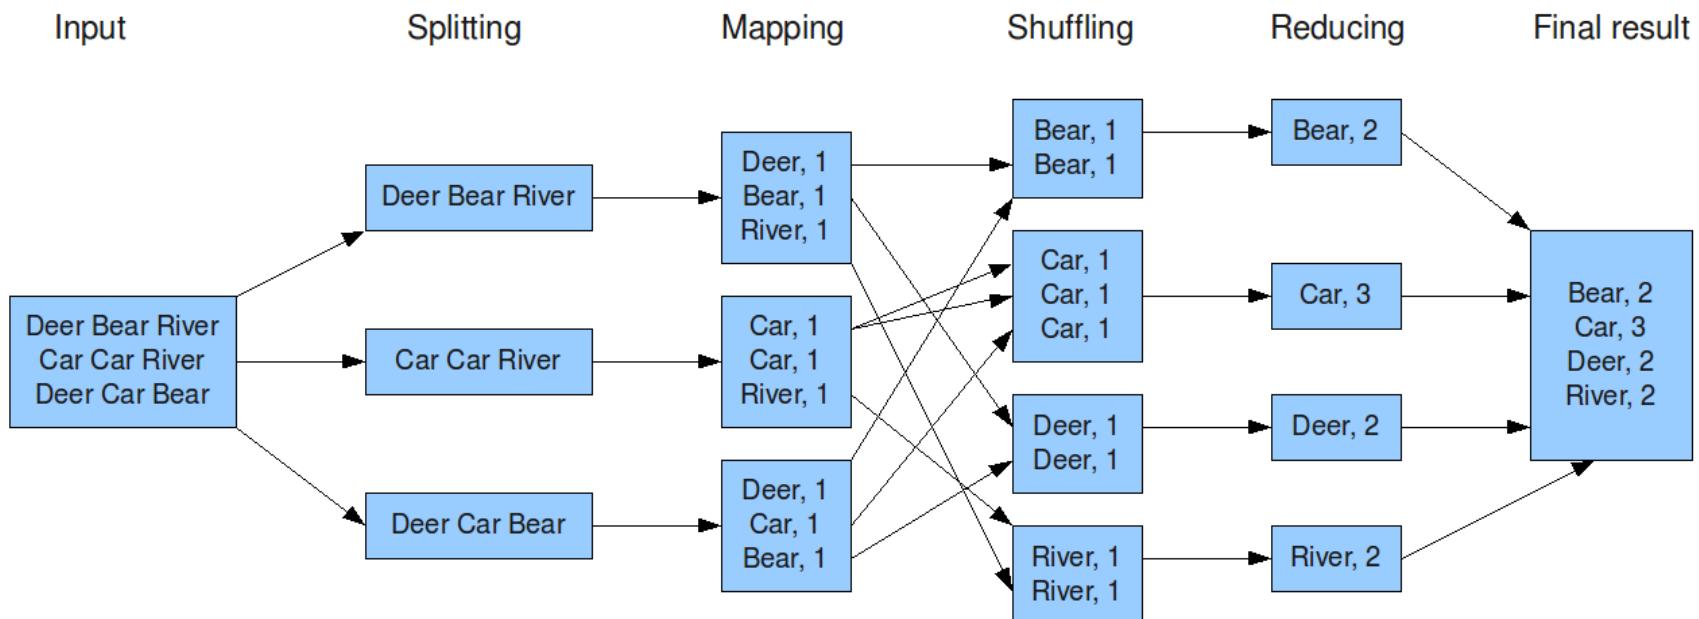


MapReduce: Google



MapReduce

The overall MapReduce word count process



shuffling : mix recuder input, sorting

MapReduce

```
import java.io.IOException;
import java.util.*;

import org.apache.hadoop.fs.Path;
import org.apache.hadoop.conf.*;
import org.apache.hadoop.io.*;
import org.apache.hadoop.mapreduce.*;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.mapreduce.lib.output.TextOutputFormat;

public class WordCount {
    public static void main(String[] args){
        }
}
```



MapReduce

```
public class Map extends Mapper<LongWritable, Text, Text, IntWritable> {  
    private final static IntWritable one = new IntWritable(1);  
    private Text word = new Text();  
  
    public void map(LongWritable key, Text value, Context context) throws  
        IOException, InterruptedException {  
        String line = value.toString();  
        StringTokenizer tokenizer = new StringTokenizer(line);  
        while (tokenizer.hasMoreTokens()) {  
            word.set(tokenizer.nextToken());  
            context.write(word, one);  
        }  
    }  
}
```



MapReduce

```
public class Reduce extends Reducer<Text, IntWritable, Text, IntWritable> {  
  
    public void reduce(Text key, Iterable<IntWritable> values, Context context)  
        throws IOException, InterruptedException {  
        int sum = 0;  
        for (IntWritable val : values) {  
            sum += val.get();  
        }  
        context.write(key, new IntWritable(sum));  
    }  
}
```



```
public class WordCount {  
    public static void main(String[] args) throws Exception {  
        Configuration conf = new Configuration();  
        Job job = new Job(conf, "wordcount");  
  
        job.setInputFormatClass(TextInputFormat.class);  
        job.setOutputFormatClass(TextOutputFormat.class);  
  
        job.setMapperClass(Map.class);  
        job.setReducerClass(Reduce.class);  
  
        job.setOutputKeyClass(Text.class);  
        job.setOutputValueClass(IntWritable.class);  
  
        FileInputFormat.addInputPath(job, new Path(args[0]));  
        FileOutputFormat.setOutputPath(job, new Path(args[1]));  
  
        job.waitForCompletion(true);  
    }  
}
```



FS Shell

Options

cat, chgrp, chmod, chown, copyFromLocal, copyToLocal, count, cp, du, dus, expunge, get, getmerge, ls, lsr, mkdir, moveFromLocal, moveToLocal, mv, put, rm, rmr, setrep, stat, tail, test, touchz

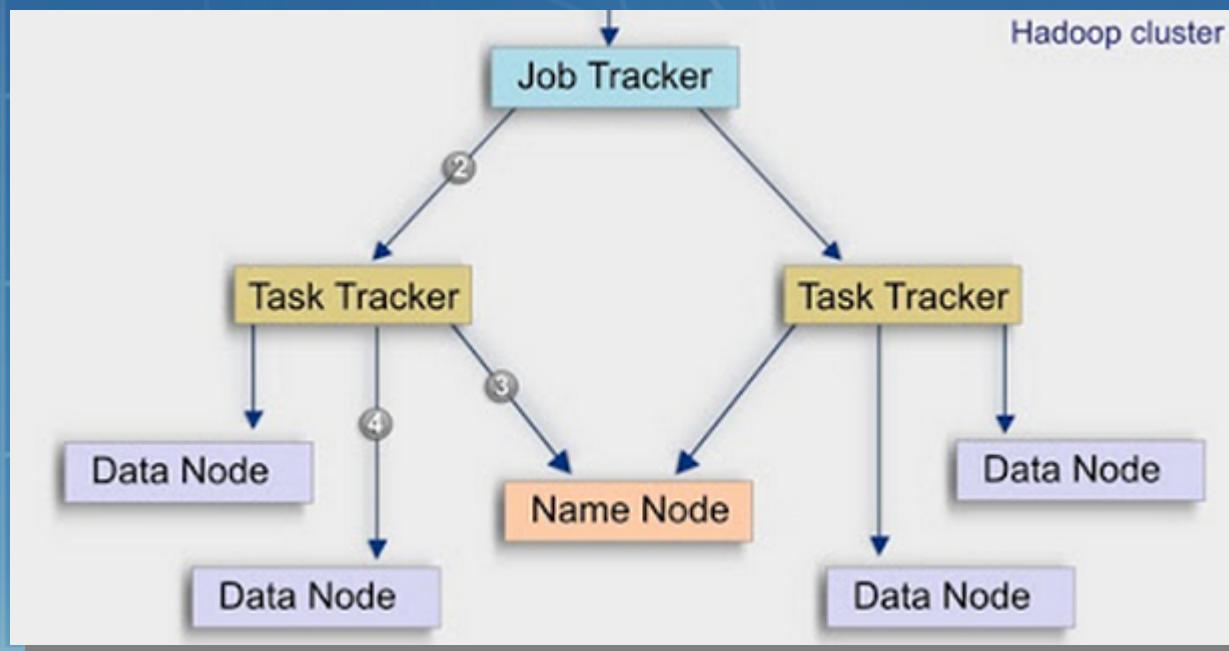
```
hadoop fs -cat /user/hadoop/file
```

```
hadoop fs -put localfile /user/hadoop/hadoopfile
```



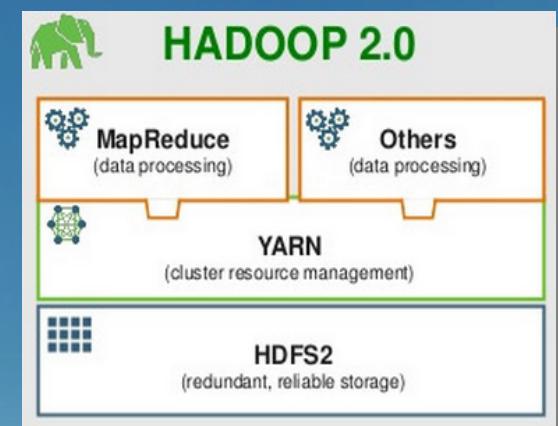
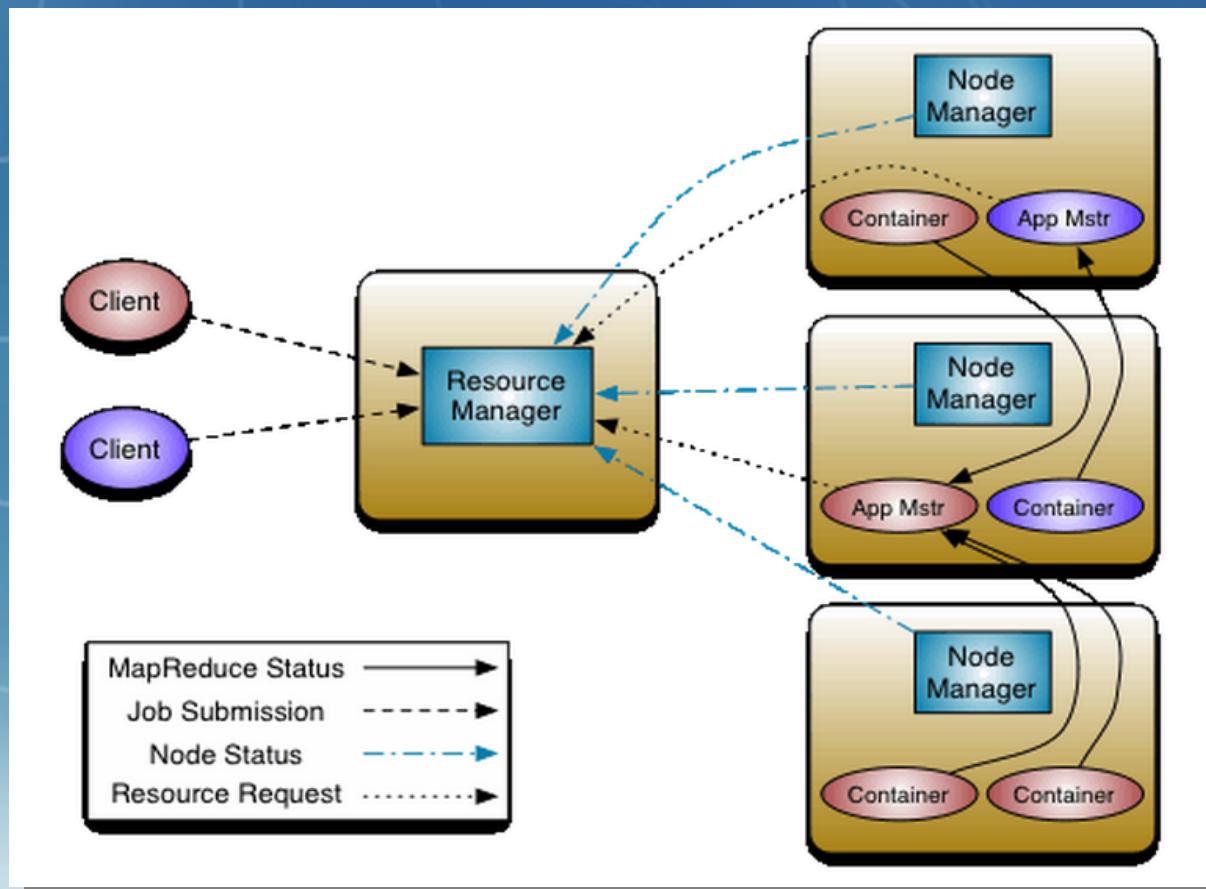
MapReduce(Hadoop 1)

- Jobtracker
- TaskTracker



NextGen MapReduce (YARN:Hadoop 2)

- Jobtracker → ResourceManager
- TaskTracker → NodeManager



- Hadoop Streamming
Utility
None java programming
Use stdin, stdout
Only Text



Python MapReduce using Hadoop Streaming
use sys.stdin to read data and sys.stdout to print data

```
$cat data | mapper.py | sort | reducer.py
```

```
cat mobydick.txt | ./mapper.py | sort | ./reducer.py >  
out/mobydick_cooccur_local.txt
```

```
head | out/mobydick_cooccur_local.tx
```

```
hadoop jar $HADOOP_HOME/hadoop-streaming.jar \  
$PARAMS \  
-mapper mapper.py \  
-reducer reducer.py \  
-input cooccur_example/mobydick.txt \  
-output cooccur_example/mobydick_out \  
-file mapper.py -file reducer.py
```



Python Streaming Mapper

```
*****  
*****mapper.py*****  
*****  
#!/usr/bin/env python  
import sys, re  
  
for line in sys.stdin:  
    # remove non alpha-numeric  
    line = re.sub("[^a-zA-Z0-9]", " ", line)  
  
    # remove leading and trailing whitespaces  
    line = line.strip()  
  
    # split the line into words  
    words = line.split()  
  
    # increase counters  
    for word in words:  
        # tab delimited; this word happened one time  
        # write the results to STDOUT (standard output);  
        # what we output here will be the input for the  
        # Reduce step, i.e. the input for reducer.py  
        #  
        # tab-delimited; the trivial word count is 1  
        print '%s\t%s' % (word, 1)
```



Python Streaming Reducer

```
*****
*****reducer.py*****
#!/usr/bin/env python
from operator import itemgetter
import sys
current_word = None
current_count = 0
word = None
for line in sys.stdin:
    # remove leading and trailing whitespace
    line = line.strip()
    # parse the mapper input
    word, count = line.split('\t', 1)
    # convert count (currently a string) to int
    try:
        count = int(count)
    except ValueError:
        # count wasn't a number, so we
        # (silently) discard the line
        continue
    # this IF-switch only works because Hadoop sorts map output
    # by key (here: word) before it is passed to the reducer
    if current_word == word:
        current_count += count
    else:
        if current_word:
            # write result to STDOUT
            print '%s\t%s' % (current_word, current_count)
        current_count = count
        current_word = word
    # do not forget to output the last word if needed!
if current_word == word:
    print '%s\t%s' % (current_word, current_count)
```



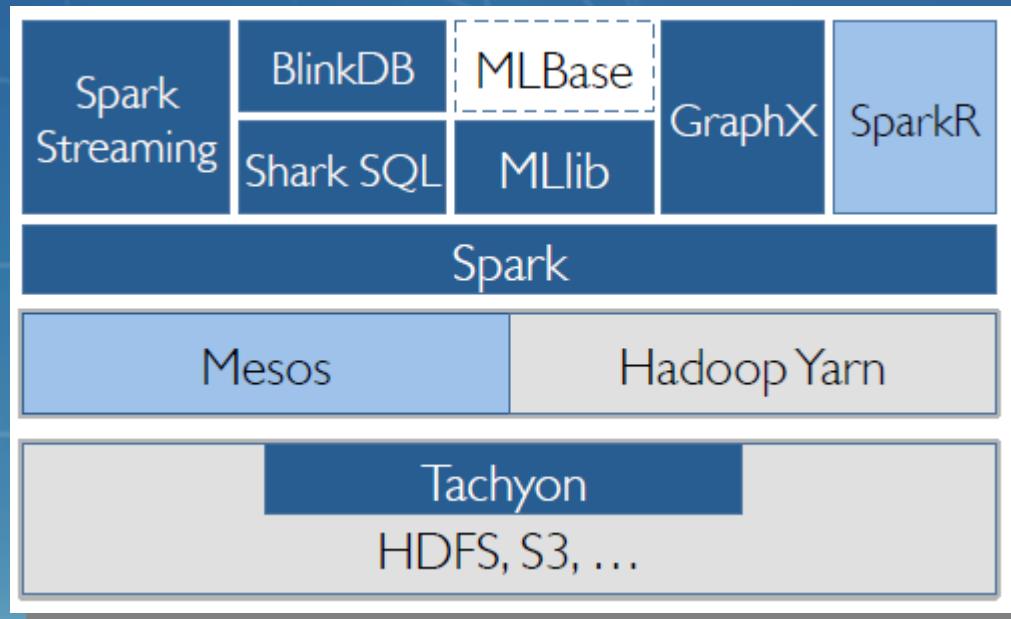
Outline

- ◆ Bigdata
- ◆ Hadoop
- ◆ **Spark**
- ◆ 3rd party
- ◆ Machine Learning
- ◆ Streaming
- ◆ Future



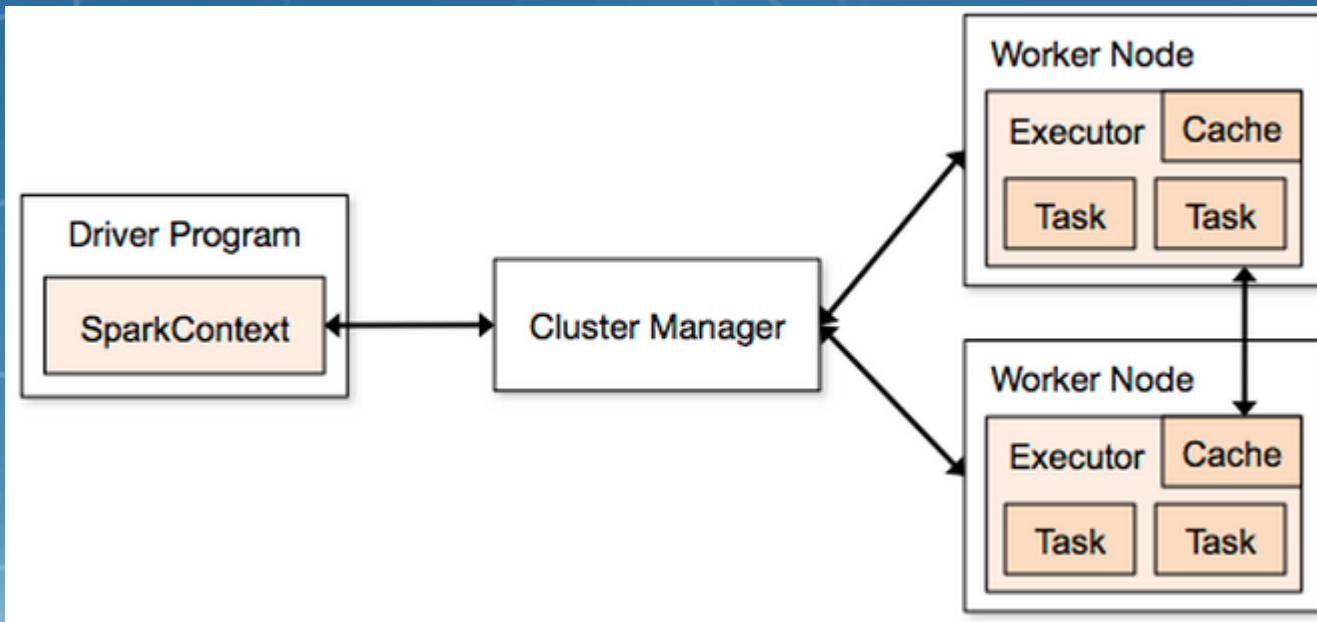
What is Spark

- A fast and general engine for large-scale data processing
- Spark 1.1.0 (latest release)

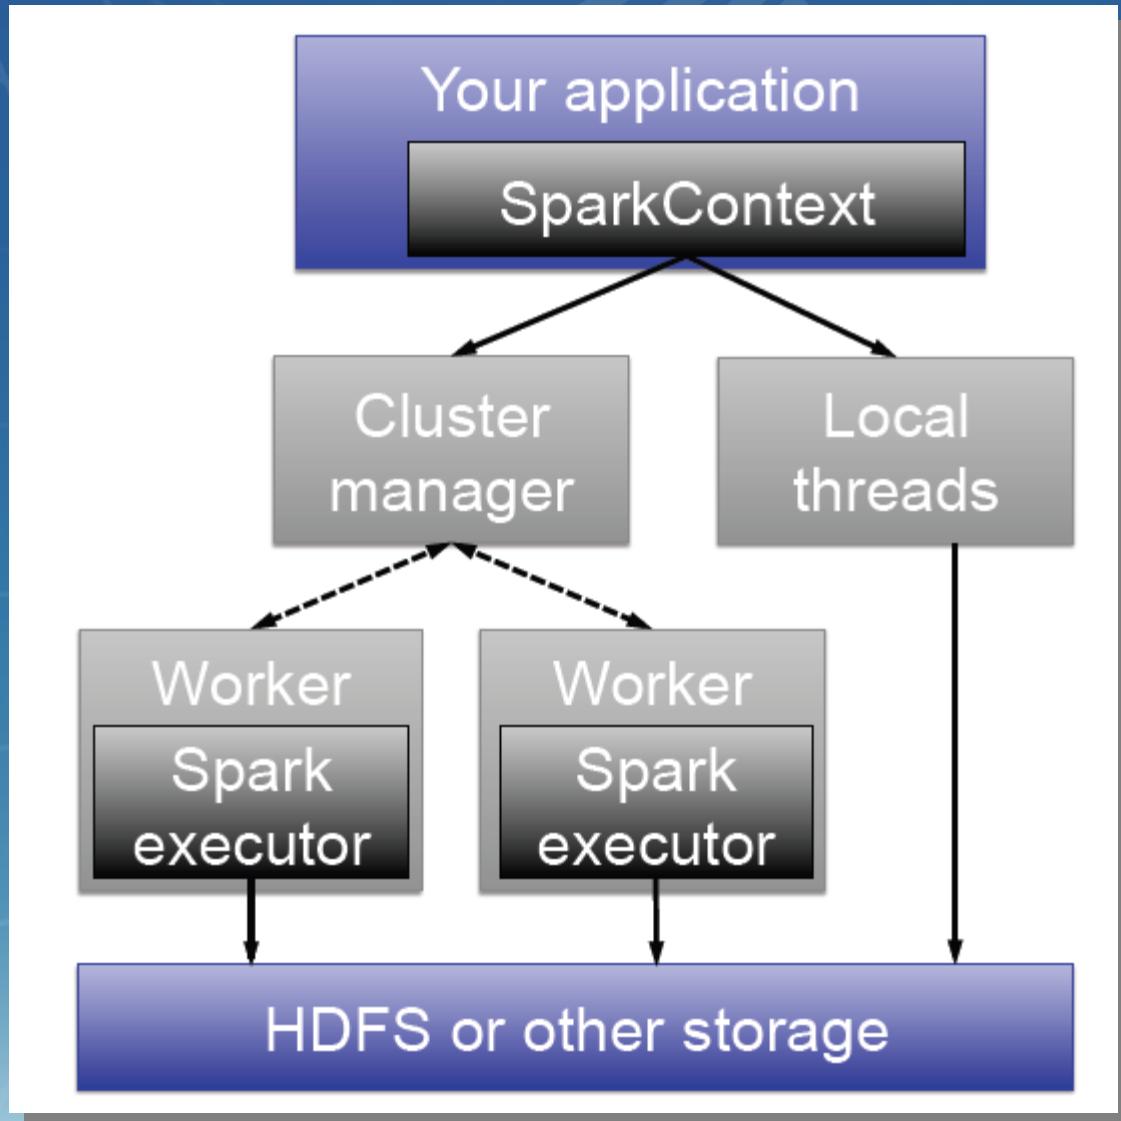


Process

```
val conf = new SparkConf().setAppName("Simple Application")
val sc = new SparkContext(conf)
val textFile = sc.textFile("README.md")
```



Components:



RDD

Resilient Distributed Datasets

- Collections of objects spread across a cluster, stored in RAM or on Disk
- Built through parallel transformations
- Automatically rebuilt on failure
- Standard RDD operations—map, countByValue, reduce, join
- Stateful operations—window, countByValueAndWindow, ...



Language Support

Python

```
lines = sc.textFile(...)  
lines.filter(lambda s: "ERROR" in s).count()
```

Scala

```
val lines = sc.textFile(...)  
lines.filter(x => x.contains("ERROR")).count()
```

Java

```
JavaRDD<String> lines = sc.textFile(...);  
lines.filter(new Function<String, Boolean>() {  
    Boolean call(String s) {  
        return s.contains("error");  
    }  
}).count();
```

Standalone Programs

- Python, Scala, & Java

Interactive Shells

- Python & Scala

Performance

- Java & Scala are faster due to static typing
- ...but Python is often fine

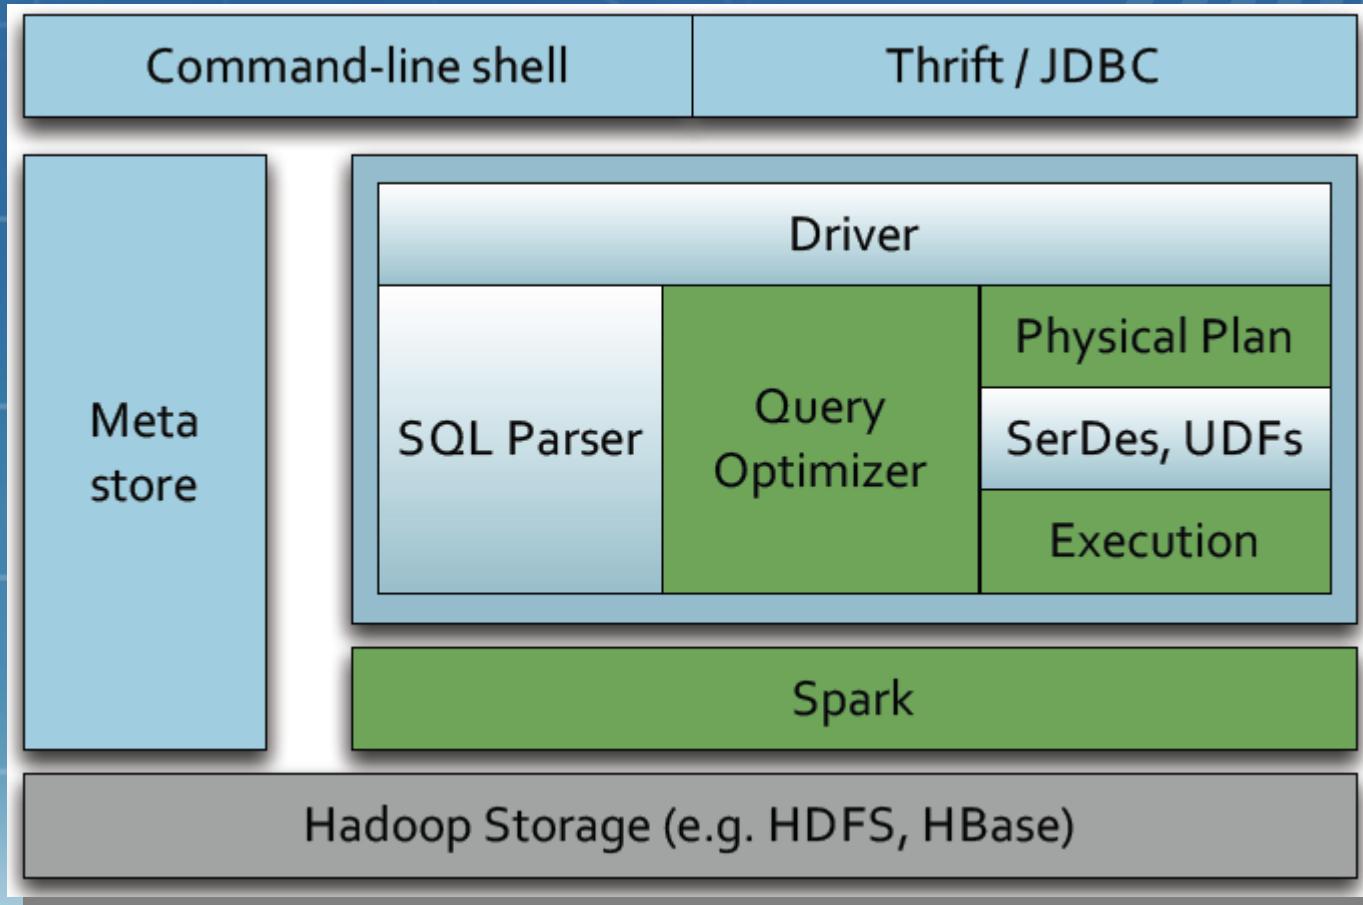
Spark SQL

- **Unifies access** to structured data
- integrated APIs in Python, Scala and Java

```
sqlCtx = new HiveContext(sc)
results = sqlCtx.sql(
    "SELECT * FROM people")
names = results.map(lambda p: p.name)
```

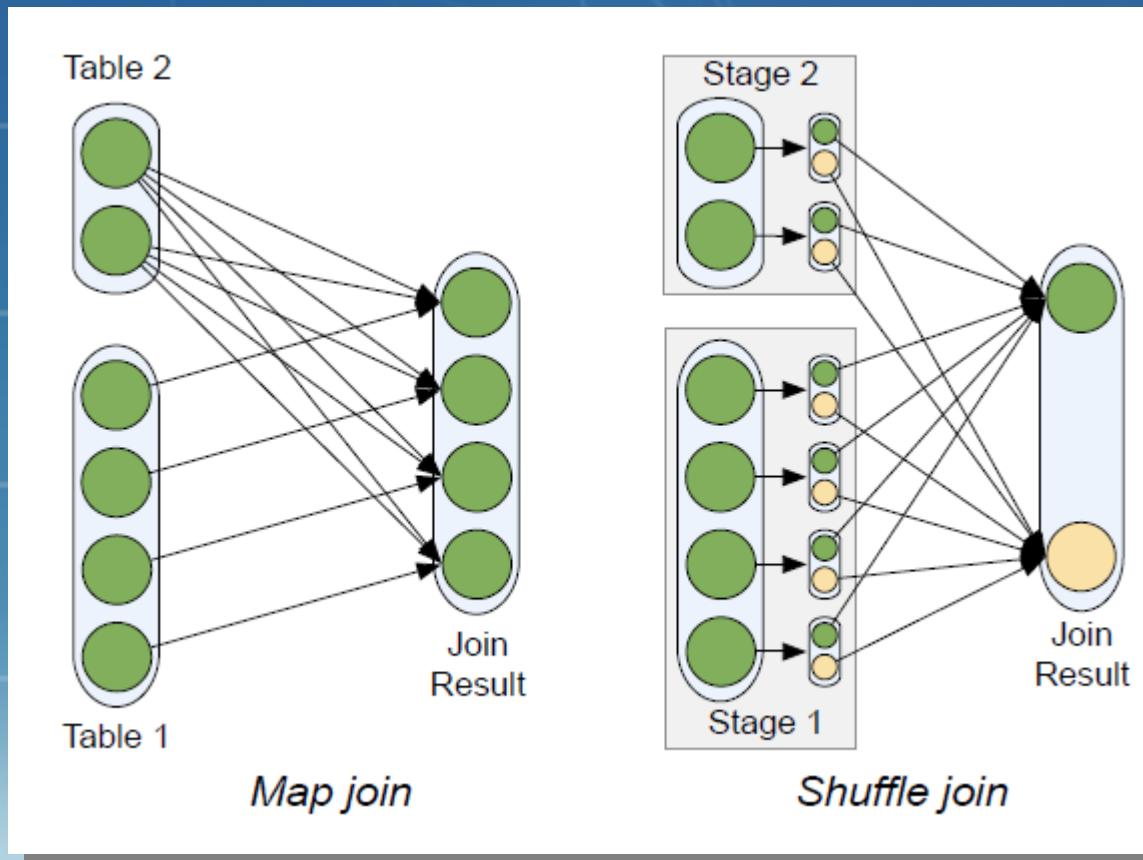


Spark SQL Shark Architecture



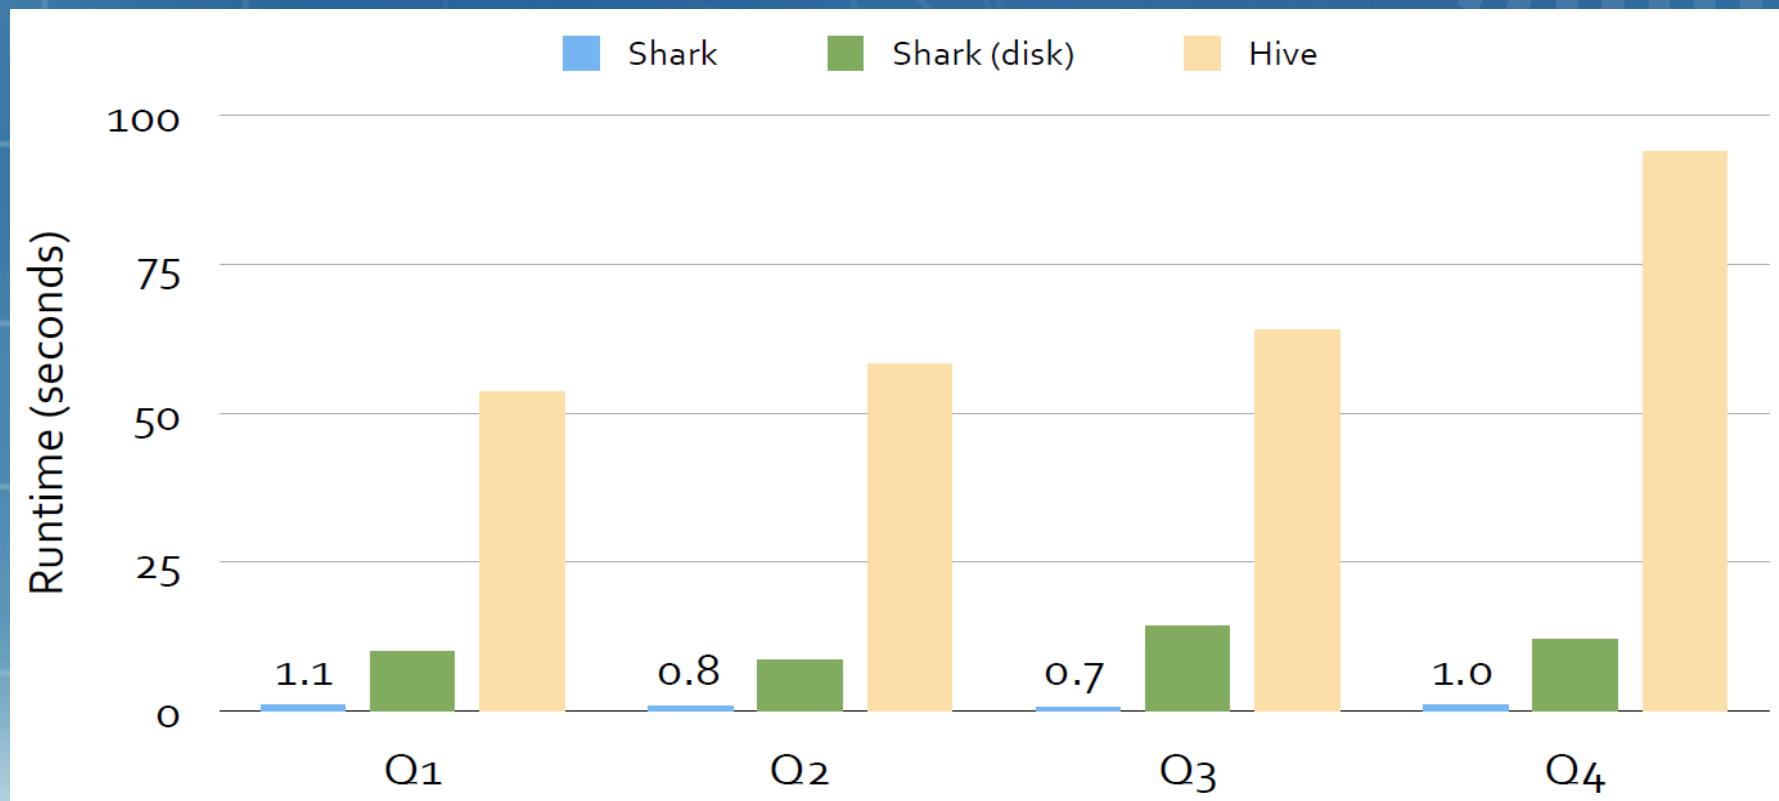
Spark SQL Shark

Partial DAG Execution (PDE)



Spark SQL Shark

Compare to hive



Spark SQL Shark

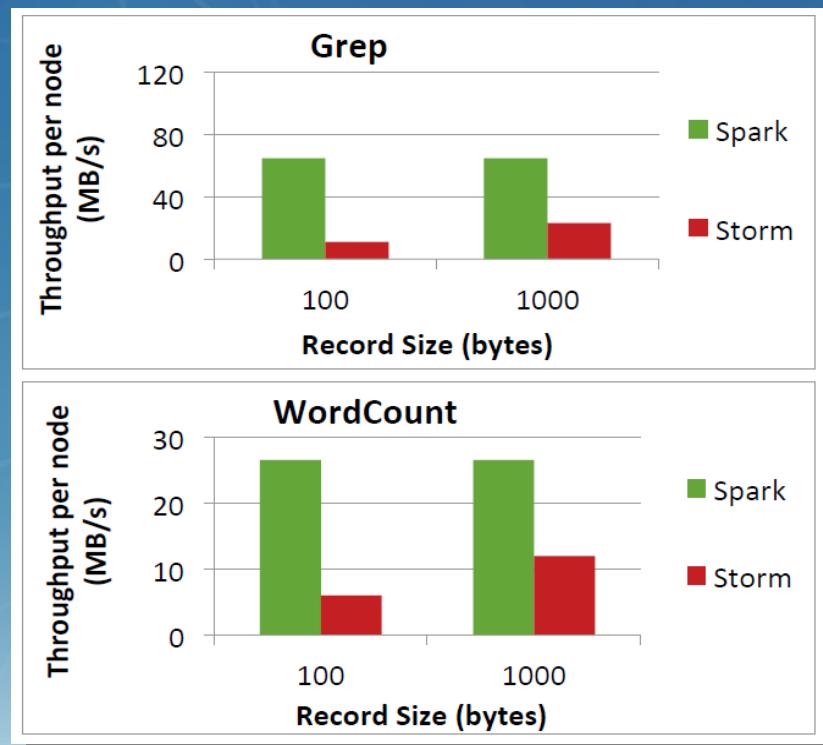
Compare to Impala

	Shark	Impala
Focus	integrate SQL with complex analytics	data warehouse / OLAP
Execution	Spark (MapReduce like)	Parallel Databases
In-memory	in-memory tables	no (buffer cache)
Fault-tolerance	tolerate slave failures	no
Large (out-of-core) joins	yes	no
UDF	yes	no



Spark Streaming

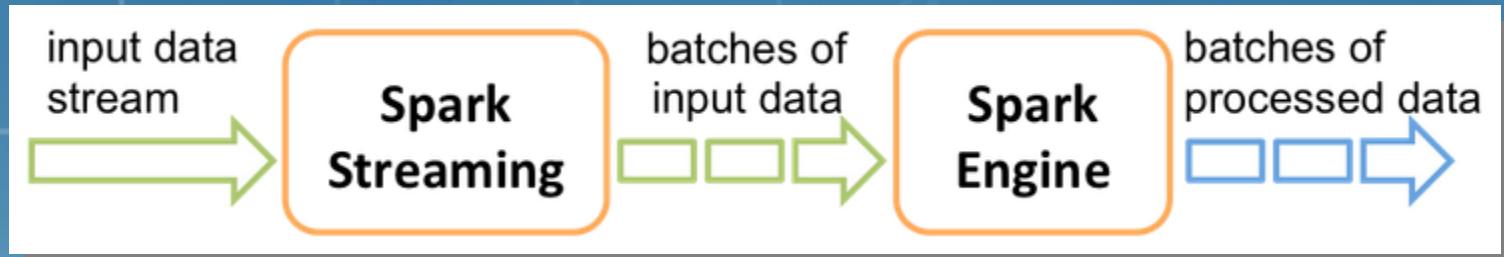
- makes it easy to build **scalable fault-tolerant streaming** applications.
- powerful **interactive** applications
- provide results in **near-real-time**



Spark Streaming

```
import org.apache.spark._  
import org.apache.spark.streaming._  
import org.apache.spark.streaming.StreamingContext._
```

```
val conf = new SparkConf().setMaster("local[2]").setAppName("NetworkWordCount")  
val ssc = new StreamingContext(conf, Seconds(1))  
val lines = ssc.socketTextStream("localhost", 9999)
```



Batch to memory : Mini batch

Spark Streaming

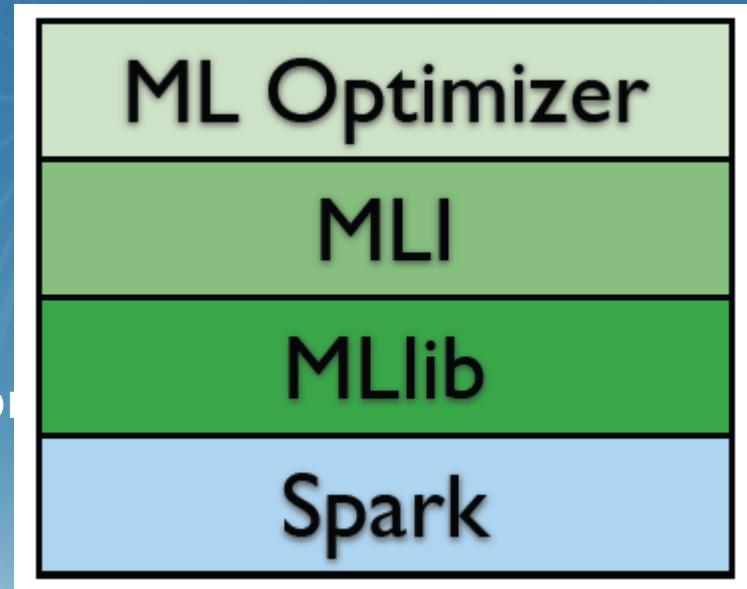
- Kafka spark-streaming-kafka_2.10
- Flume spark-streaming-flume_2.10
- Kinesis spark-streaming-kinesis-asl_2.10 [Apache Software License]
- Twitter spark-streaming-twitter_2.10
- ZeroMQ spark-streaming-zeromq_2.10
- MQTT spark-streaming-mqtt_2.10



MLlib

scalable machine learning library Usable in Java, Scala and Python, 100x faster than MapReduce.

- linear models (SVMs, logistic regression, linear regression)
- decision trees
- naive Bayes
- alternating least squares (ALS)
- k-means
- singular value decomposition (SVD)
- principal component analysis (PCA)
- Feature extraction and transformation
- stochastic gradient descent
- limited-memory BFGS (L-BFGS)



GraphX

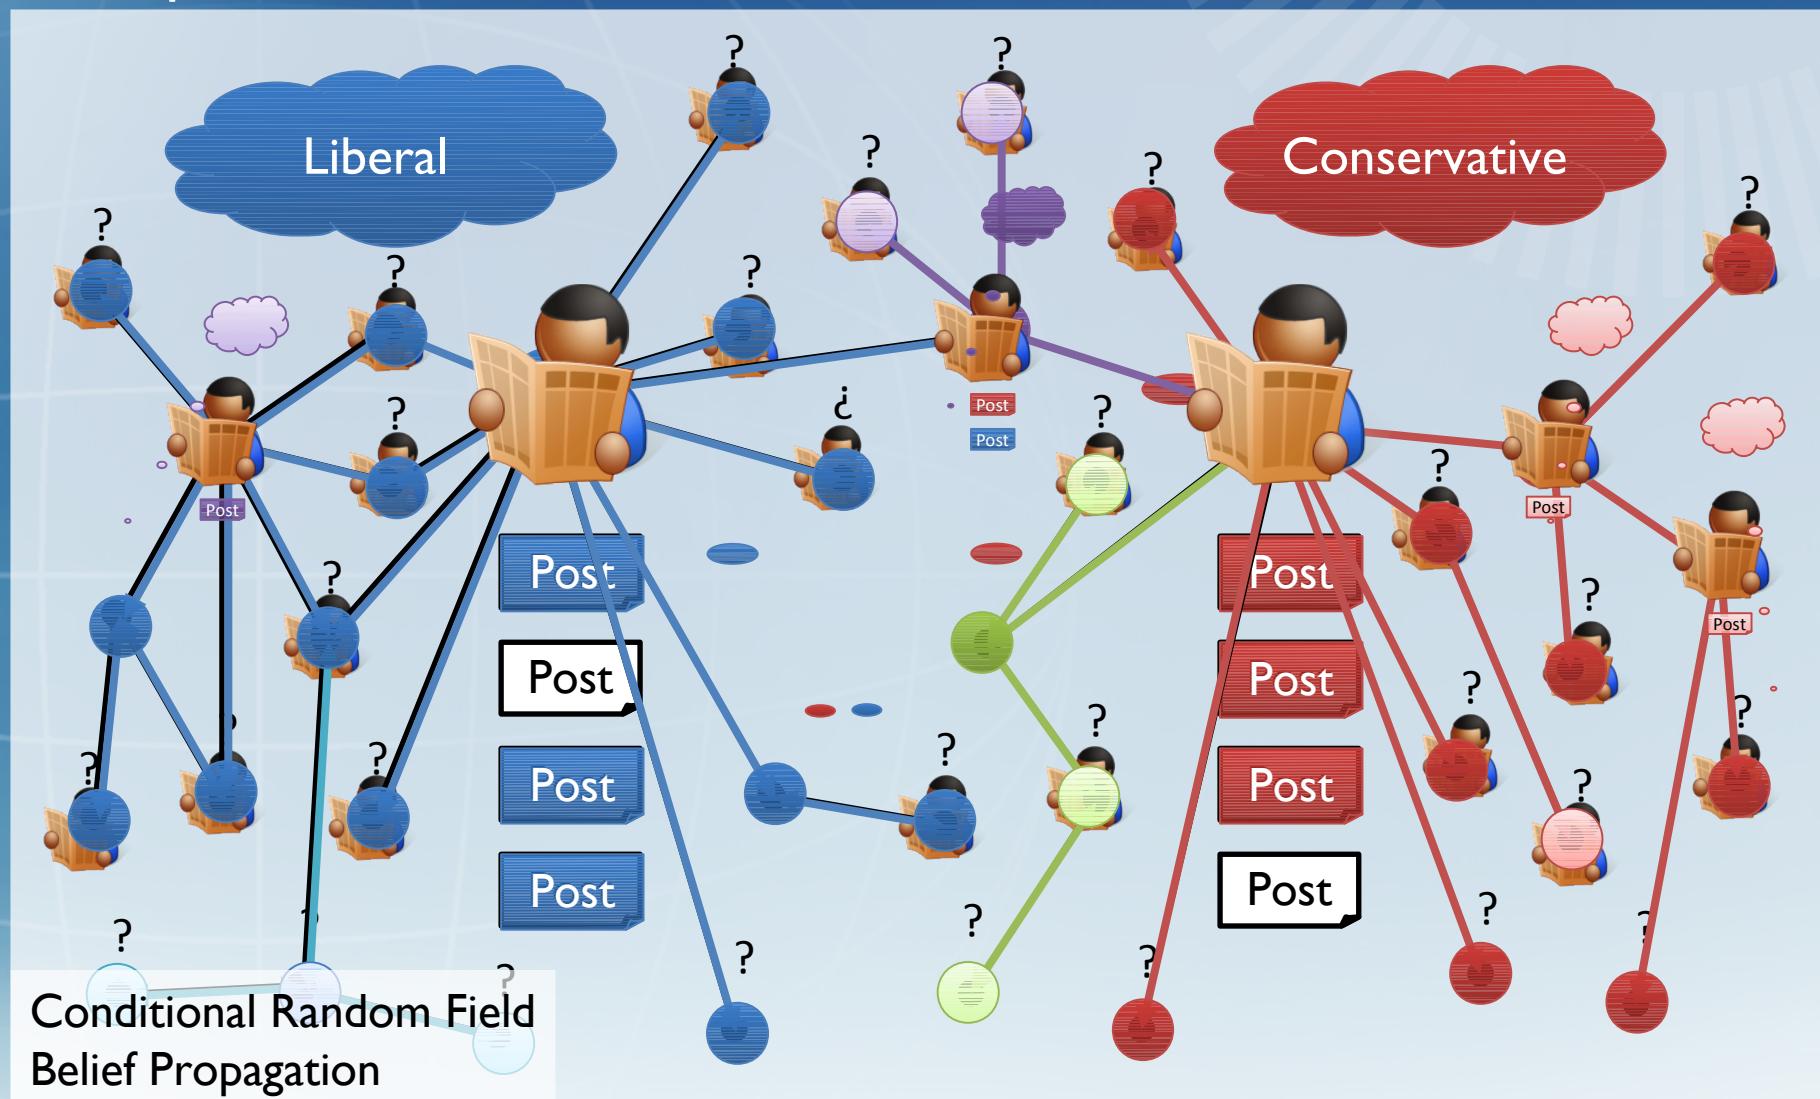
API for graphs and graph-parallel computation.

- Graph-parallel primitives on Spark.
- Currently slower than GraphLab, but
 - No need for specialized systems
- Easier ETL, and easier consumption of output

Interactive graph data mining

- Future work will bring performance closer to specialized engines.





GraphX

More Graph algorithms

Collaborative Filtering

Alternating Least Squares

Stochastic Gradient Descent

Tensor Factorization

SVD

Structured Prediction

Loopy Belief Propagation

Max-Product Linear Programs

Gibbs Sampling

Semi-supervised ML

Graph SSL

CoEM

Graph Analytics

PageRank

Single Source Shortest Path

Triangle-Counting

Graph Coloring

K-core Decomposition

Personalized PageRank

Classification

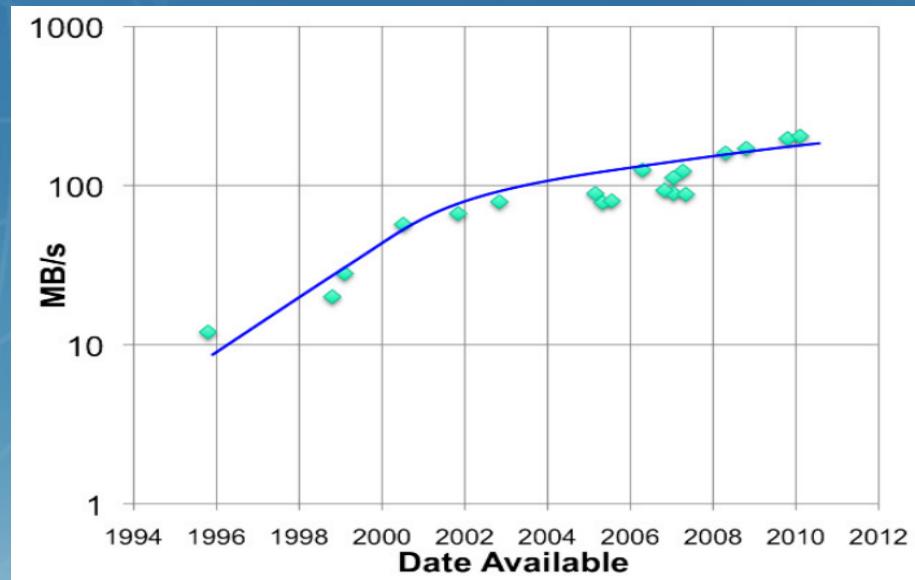
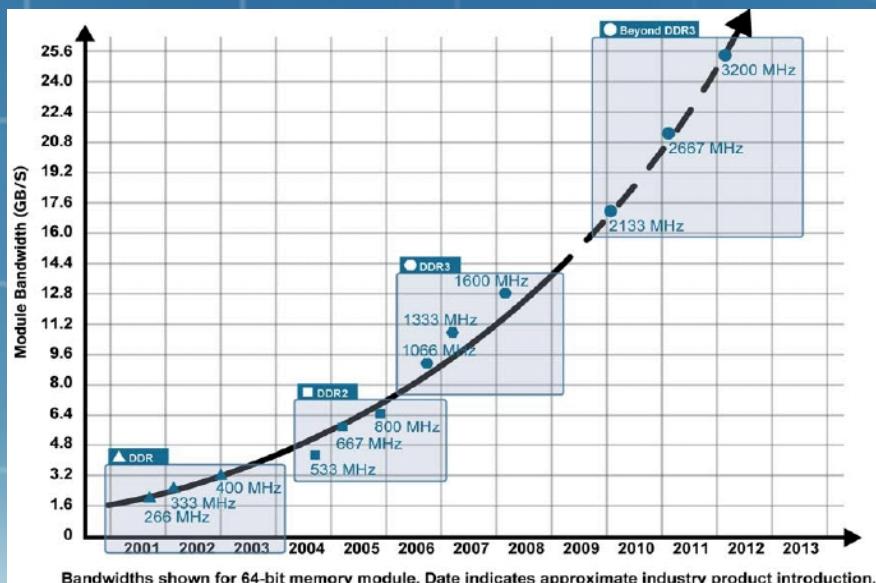
Neural Networks

Lasso



Tachyon

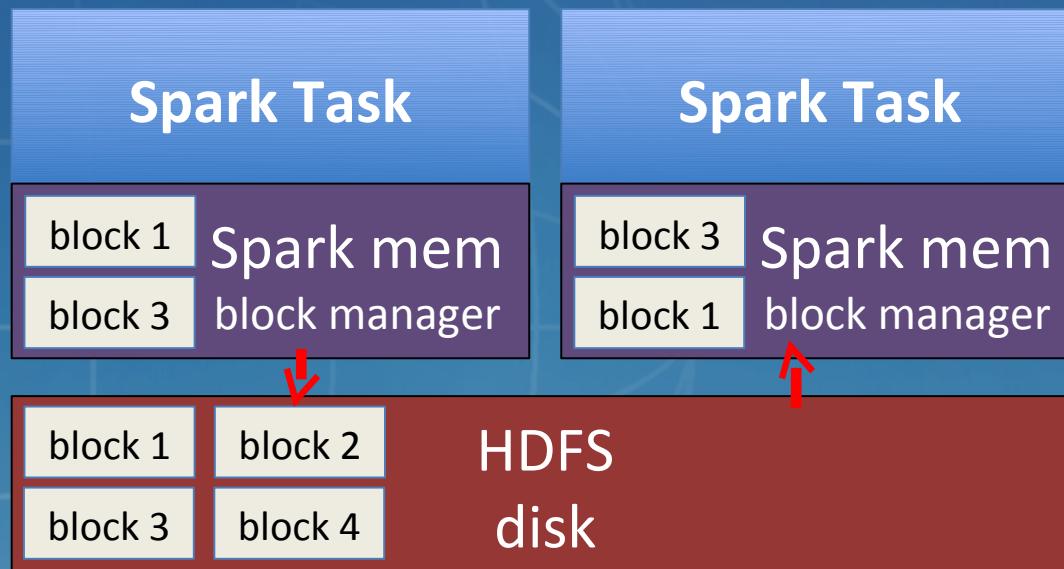
- High-throughput, fault-tolerant in-memory storage
- Interface compatible to HDFS
- Further improve performance for Spark, and Hadoop
- Growing community with 10+ organizations contributing



RAM throughput increasing exponentially, Disk slowly



Tachyon
Motive
Different jobs share

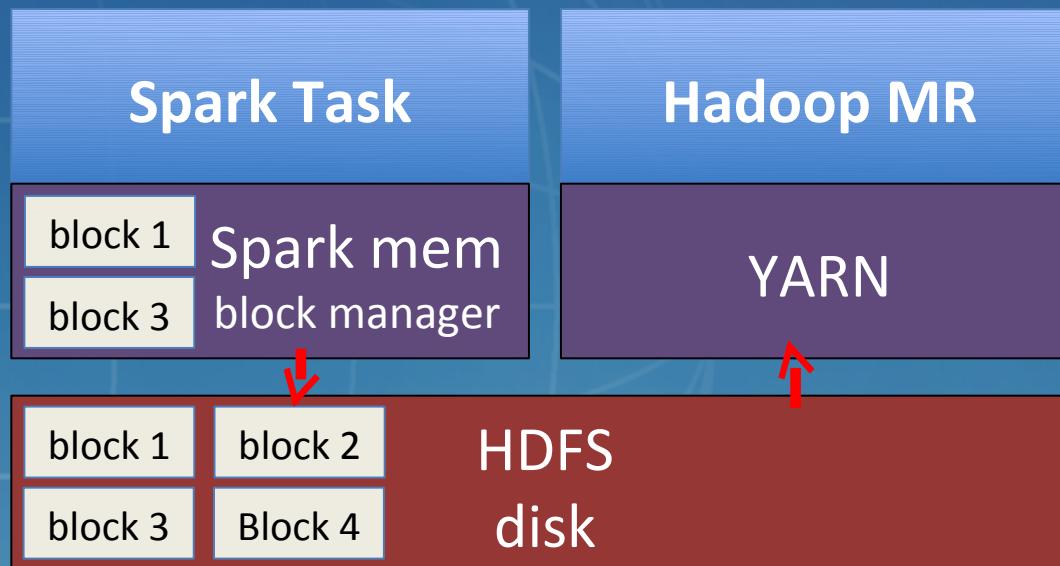


→ Slow writes to disk

Tachyon

Motive

Different framework share

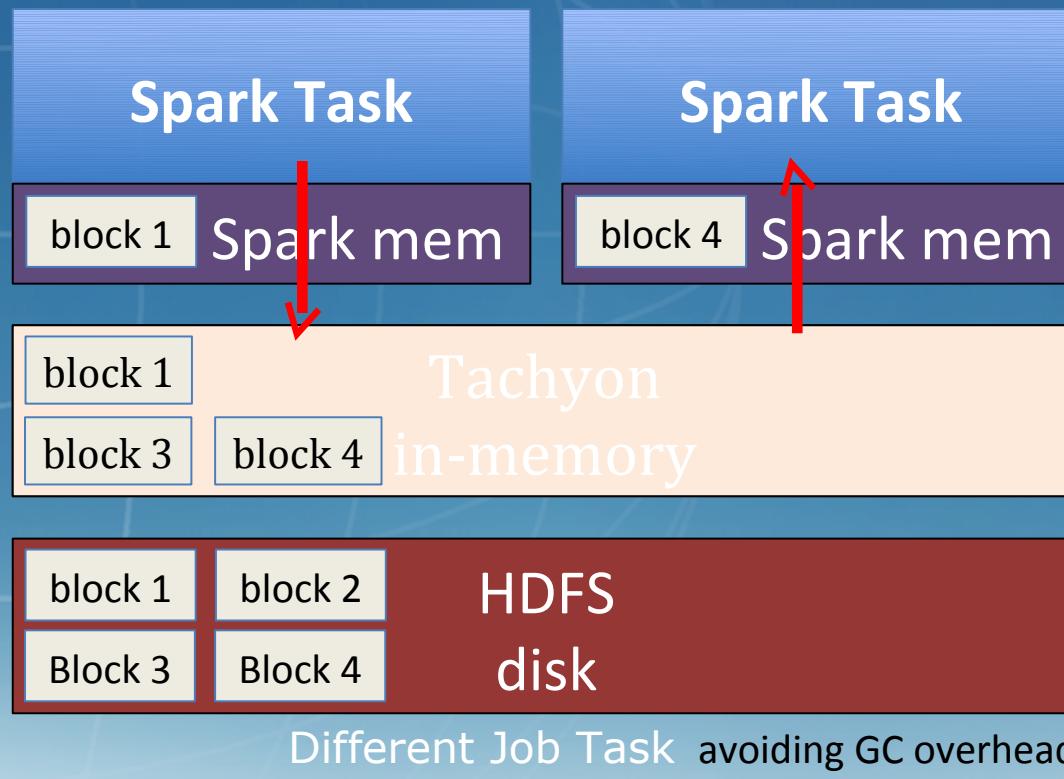


→ Slow writes to disk

Tachyon

Solution

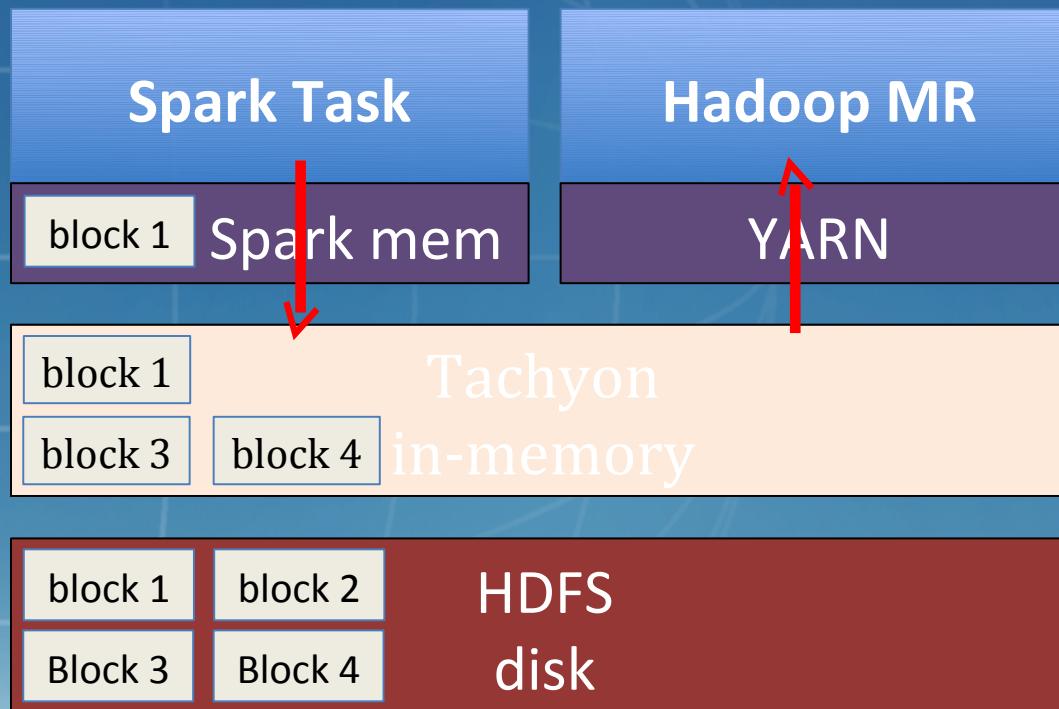
execution engine & storage engine same JVM process
(no GC & duplication)



Tachyon

Solution

execution engine & storage engine same JVM process
(no GC & duplication)

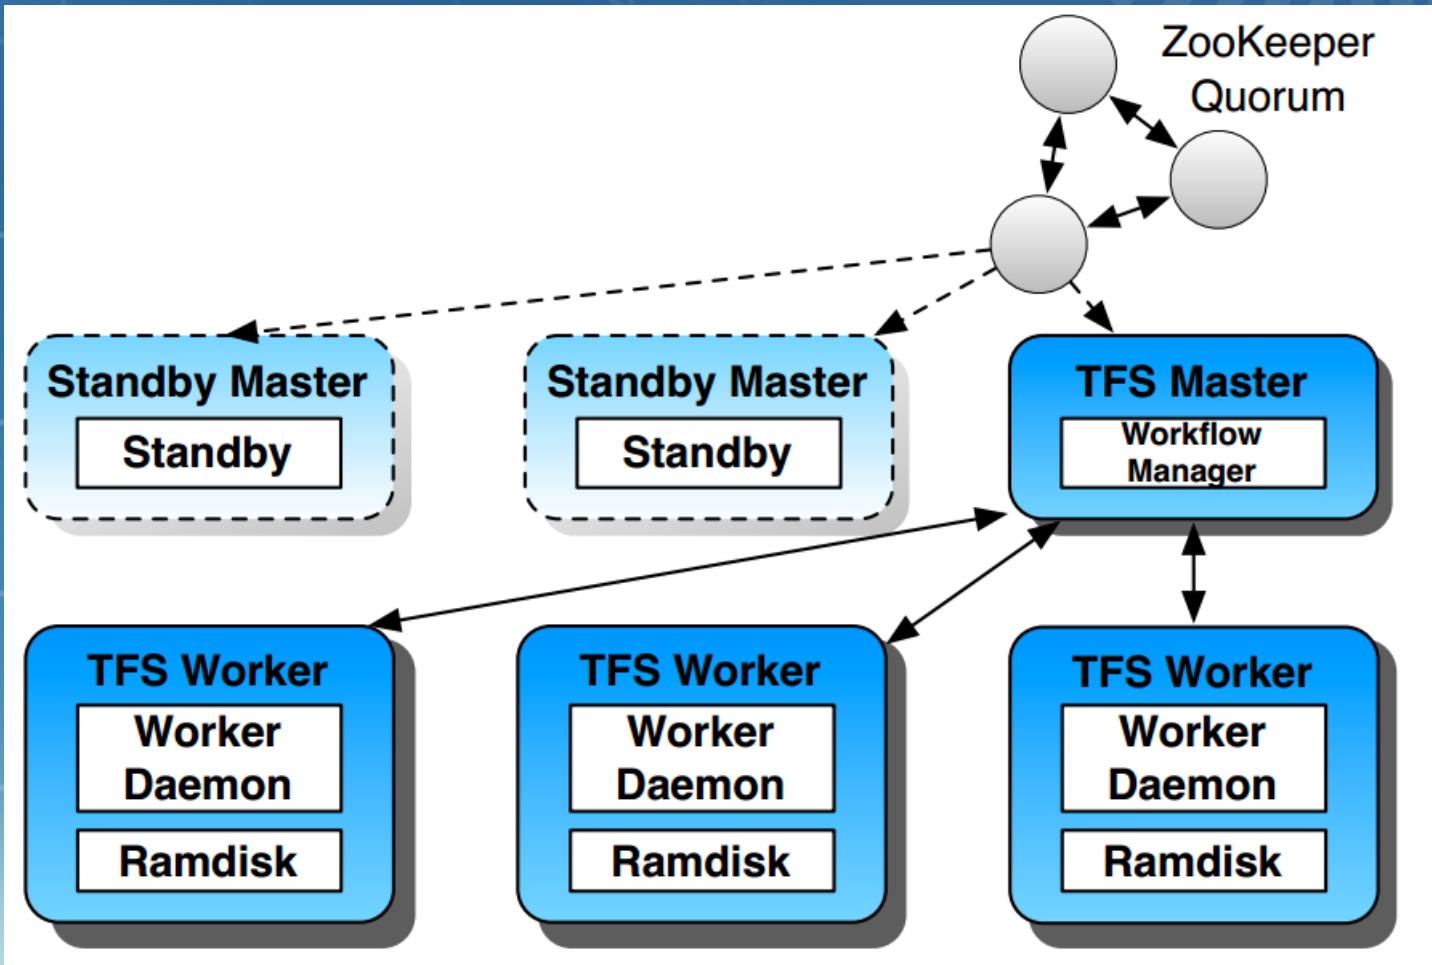


Different Framework avoiding GC overhead



Tachyon

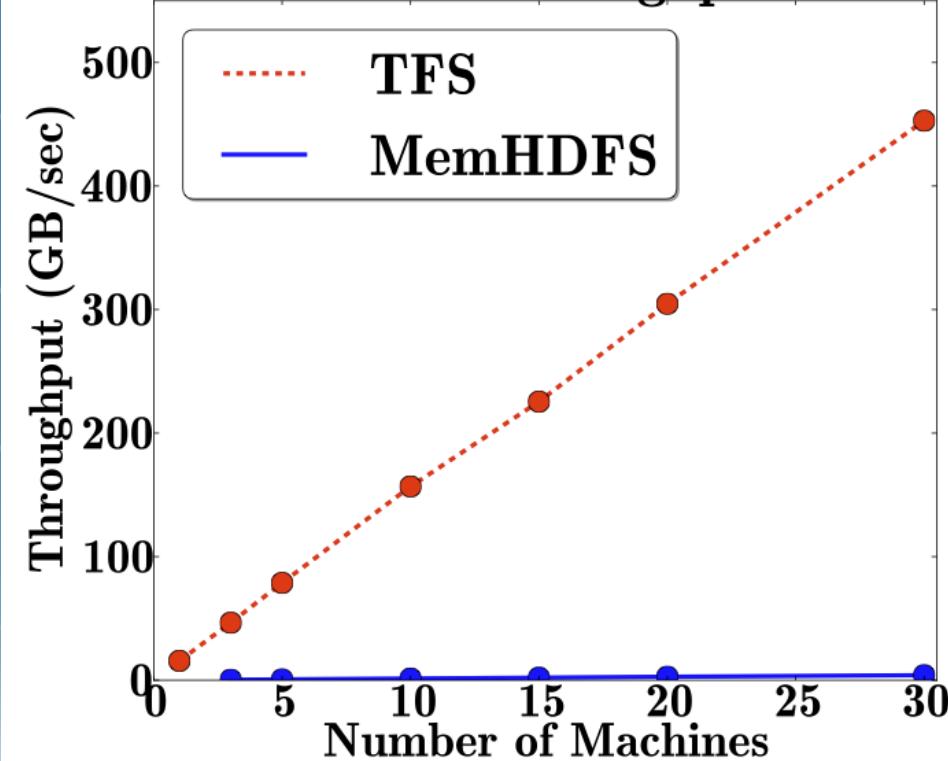
Architecture



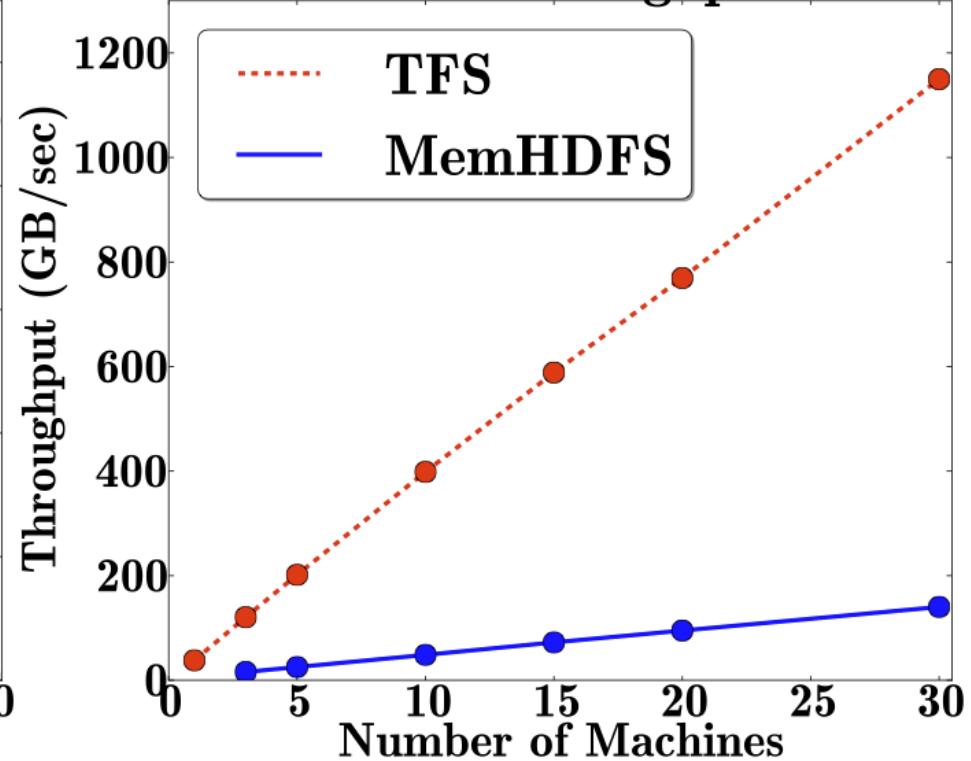
Tachyon

Comparison with in Memory HDFS

Write Throughput



Read Throughput



Tachyon

Open Source Status

- V0.4.0 (Feb 2014)
- 20 Developers (7 from Berkeley, 13 from outside)
 - 11 Companies
- Writes go synchronously to under filesystem
(No lineage information in Developer Preview release)
- MapReduce and Spark can run without any code change
(ser/de becomes the new bottleneck)



Tachyon

Using : hdfs → tachyon

- **Spark**

```
val file = sc.textFile("tachyon://ip:port/path")
```

- **Shark(Spark SQL)**

```
CREATE TABLE orders_tachyon AS SELECT * FROM  
orders;
```

- **Hadoop MapReduce**

```
hadoop jar examples.jar wordcount  
tachyon://localhost/input  
tachyon://localhost/output
```



Tachyon

**Tachyon
is in Fedora 20**

Thanks to Redhat!



BlinkDB

- A data analysis (warehouse) system
- is compatible with Apache Hive and Facebook's Presto (storage, serdes, UDFs, types, metadata)

Support **interactive** SQL-like aggregate queries over **massive sets of data**

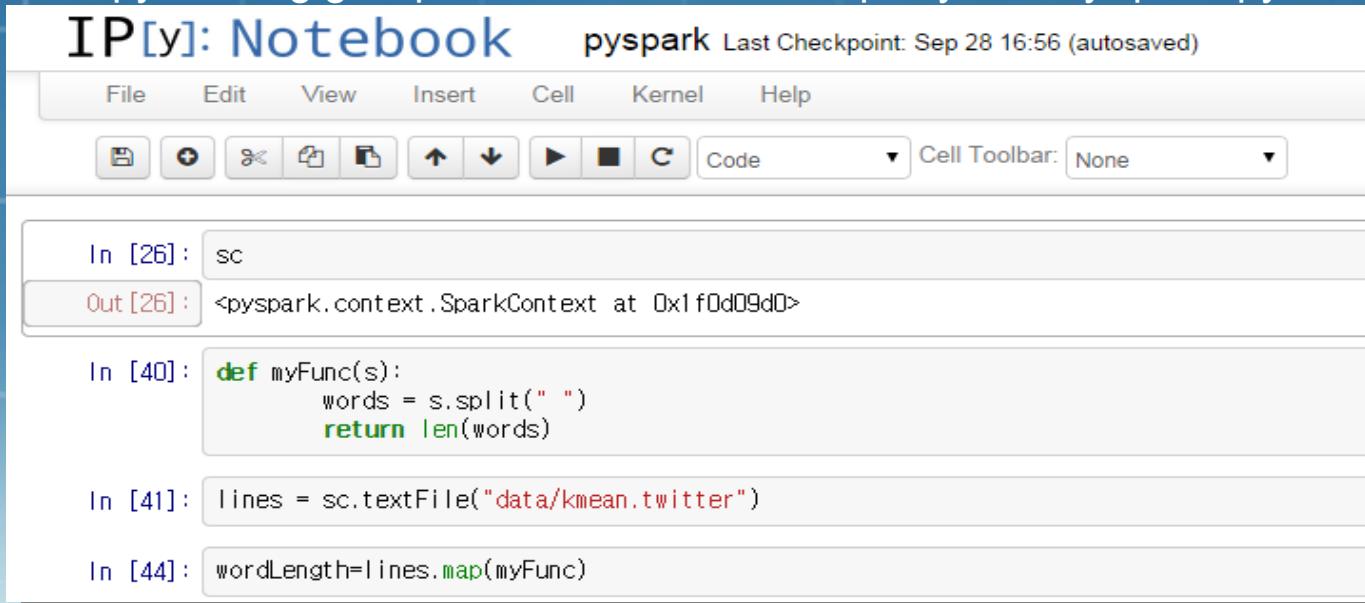


PySpark

- The Spark Python API
- All of PySpark's library dependencies, including **Py4J**, are bundled with PySpark and automatically imported.
- Can run with Ipython notebook

PySpark and Ipython nodebook setting manual link:

<http://nbviewer.ipython.org/gist/fperez/6384491/00-Setup-IPython-PySpark.ipynb>



The screenshot shows an IPython Notebook interface with the title "IP[y]: Notebook" and a status bar indicating "pyspark Last Checkpoint: Sep 28 16:56 (autosaved)". The toolbar includes buttons for file operations, cell selection, and execution. Below the toolbar, the notebook displays several code cells:

- In [26]: `sc`
Out [26]: `<pyspark.context.SparkContext at 0x1f0d09d0>`
- In [40]:

```
def myFunc(s):
    words = s.split(" ")
    return len(words)
```
- In [41]: `lines = sc.textFile("data/kmean.twitter")`
- In [44]: `wordLength=lines.map(myFunc)`



SparkR

```
install.packages("rJava")
library(devtools)
install_github("amplab-extras/SparkR-pkg", subdir="pkg")
library(SparkR)
sc <- sparkR.init(master="local")

./sparkR
```

<http://amplab-extras.github.io/SparkR-pkg/>

<http://amplab-extras.github.io/SparkR-pkg/>



Outline

- ◆ Bigdata
- ◆ Hadoop
- ◆ Spark
- ◆ **3rd Party Biobigdata**
- ◆ Machine Learning
- ◆ Streaming
- ◆ Future



Dumbo

- Python module that makes writing and running Hadoop Streaming programs very easy

<http://klbostee.github.io/dumbo/>

<https://github.com/klbostee/dumbo/wiki>

Happy

Hadoop + Python = Happy

A framework for writing map-reduce programs for
Hadoop using **Jython**

<https://code.google.com/p/happy/>



Happy

```
import sys, happy, happy.log
happy.log.setLevel("debug")
log = happy.log.getLogger("wordcount")
class WordCount(happy.HappyJob):
    def __init__(self, inputpath, outputpath):
        happy.HappyJob.__init__(self)
        self.inputpaths = inputpath
        self.outputpath = outputpath
        self.inputformat = "text"
    def map(self, records, task):
        for _, value in records:
            for word in value.split():
                task.collect(word, "1")
    def reduce(self, key, values, task):
        count = 0;
        for _ in values: count += 1
        task.collect(key, str(count))
        log.debug(key + ":" + str(count))
        happy.results["words"] = happy.results.setdefault("words", 0) + count
        happy.results["unique"] = happy.results.setdefault("unique", 0) + 1
    if __name__ == "__main__":
        if len(sys.argv) < 3:
            print "Usage: <inputpath> <outputpath>"
            sys.exit(-1)
        wc = WordCount(sys.argv[1], sys.argv[2])
        results = wc.run()
        print str(sum(results["words"])) + " total words"
        print str(sum(results["unique"])) + " unique words"
```



Octopy : Google?

```
### triangular.py
# server
source = dict(zip(range(100), range(100)))

def final(key, value):
    print key, value

# client
def mapfn(key, value):
    for i in range(value + 1):
        yield key, i

def reducefn(key, value):
    return sum(value)
```

<https://code.google.com/p/octopy/>

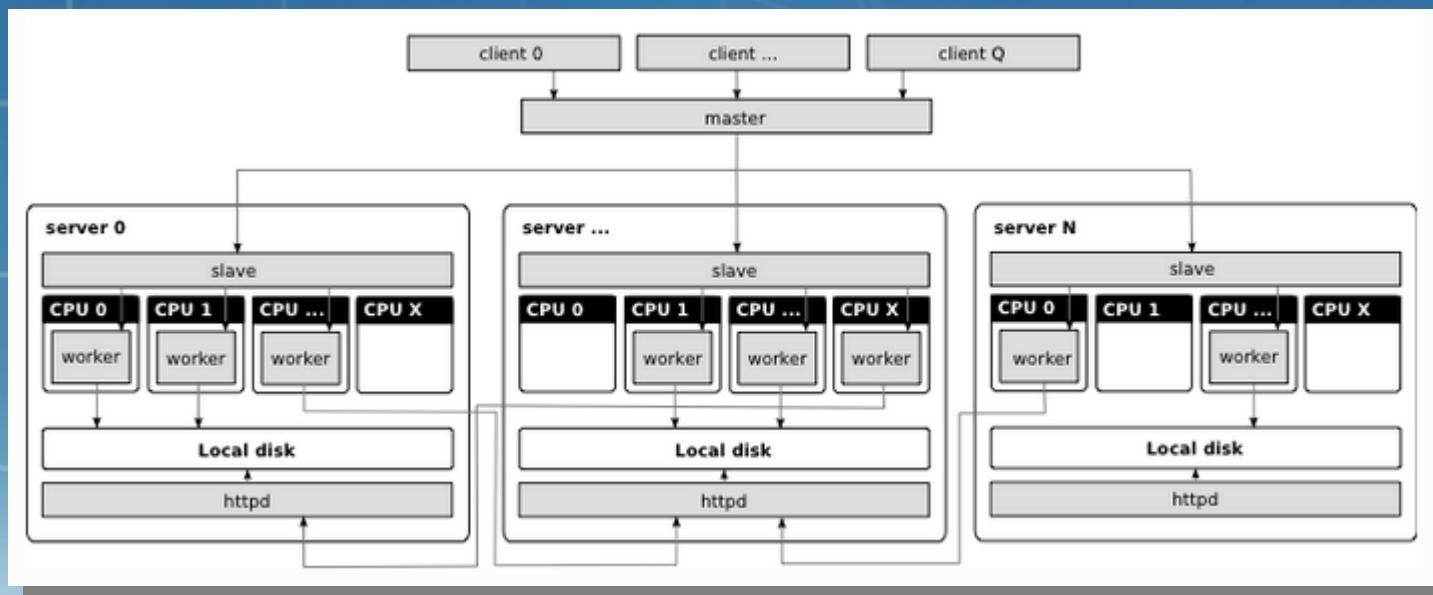


Disco

**open-source framework for distributed computing based
on the MapReduce paradigm**

Disco Distributed Filesystem

<http://discoproject.org/>





Pydoop

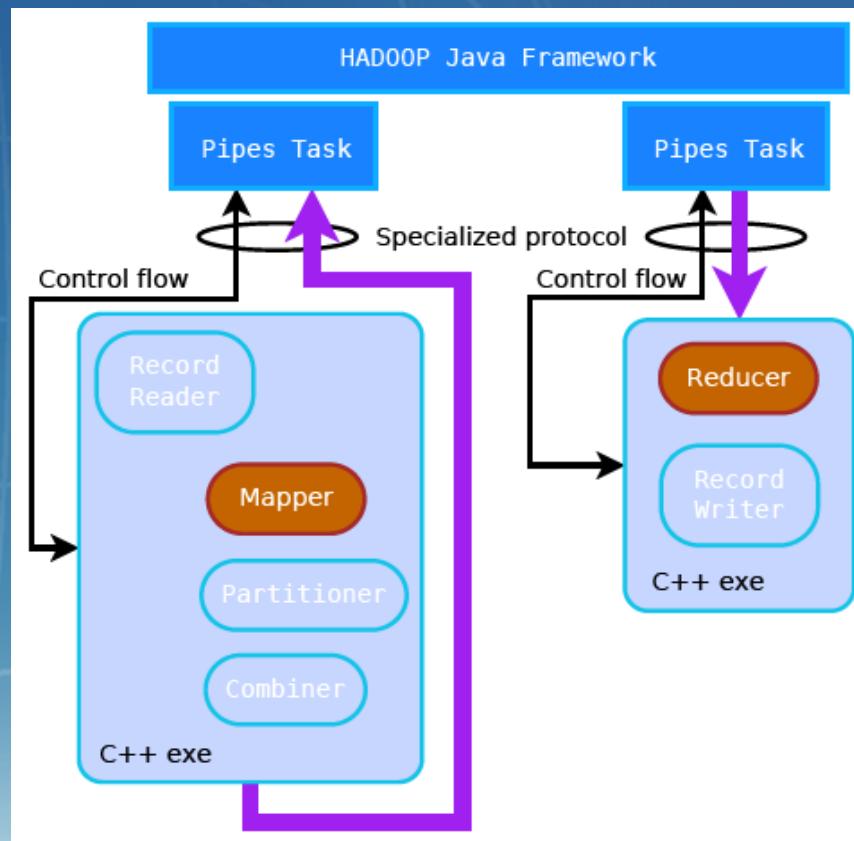
- Python API for Hadoop
- Access to most MR components, including RecordReader, RecordWriter and Partitioner
- Get configuration, set counters and report status via context objects

<http://pydoop.sourceforge.net/docs/index.html>



Pydoop

Using C/C++ APIs for both MR and HDFS are supported by Hadoop Pipes



Pydoop

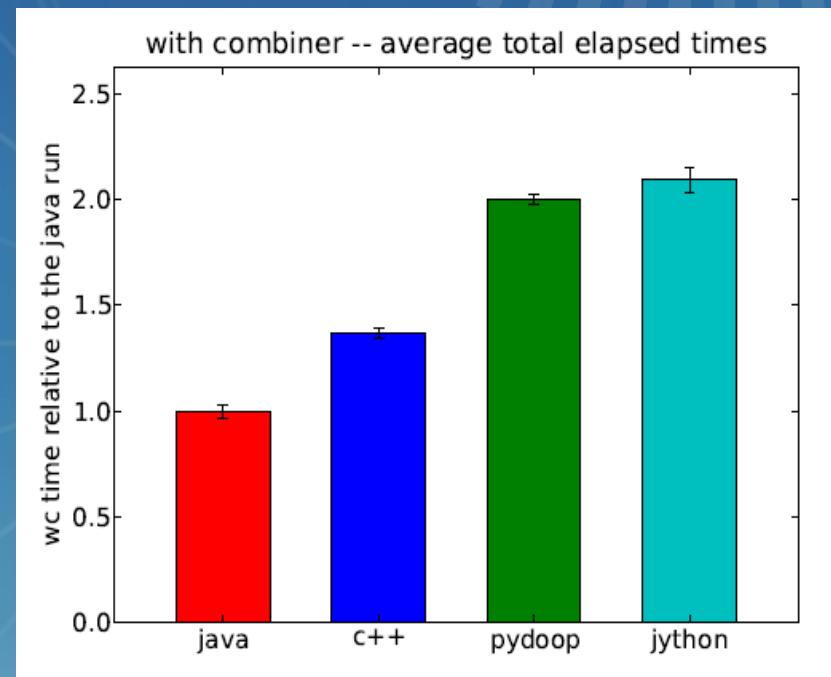
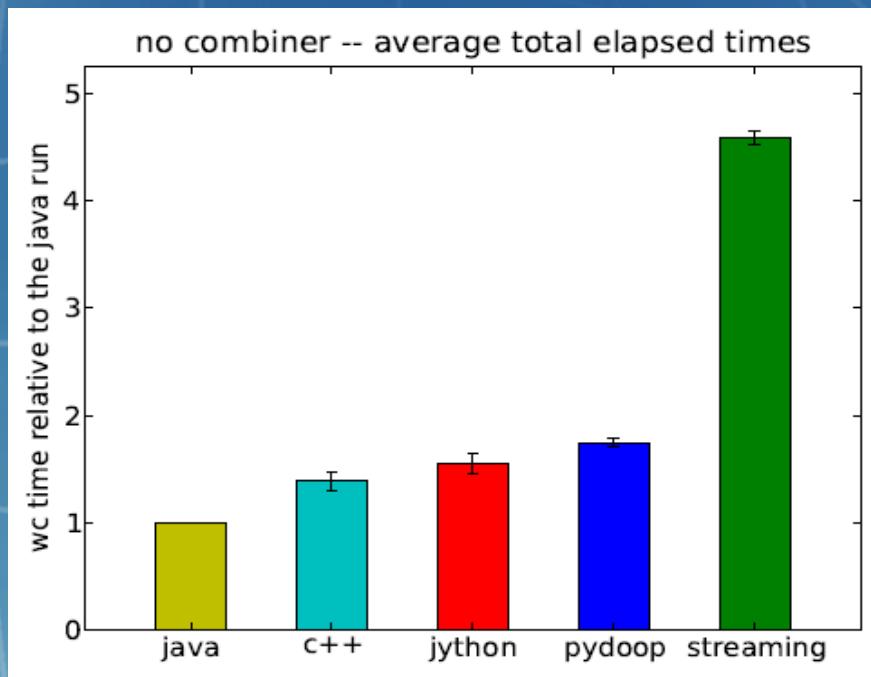
```
from pydoop.pipes import Mapper, Reducer, Factory, runTask
class WordCountMapper(Mapper):
    def map(self, context):
        words = context.getInputValue().split()
        for w in words:
            context.emit(w, "1")

class WordCountReducer(Reducer):
    def reduce(self, context):
        s = 0
        while context.nextValue():
            s += int(context.getInputValue())
        context.emit(context.getInputKey(), str(s))
runTask(Factory(WordCountMapper, WordCountReducer))
```

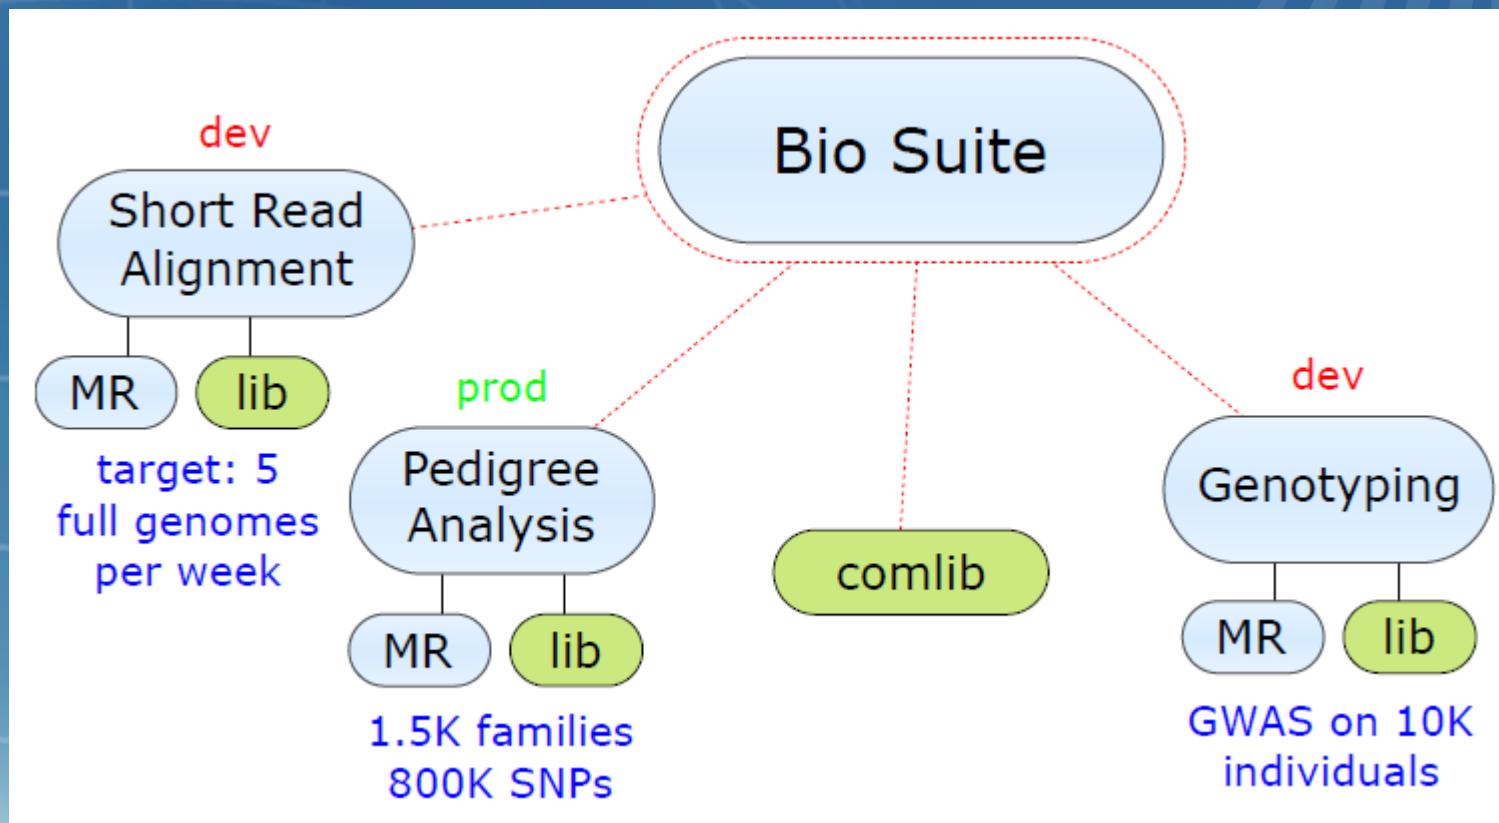


Pydoop

– Performance:



Pydoop Status





Hadoop-BAM

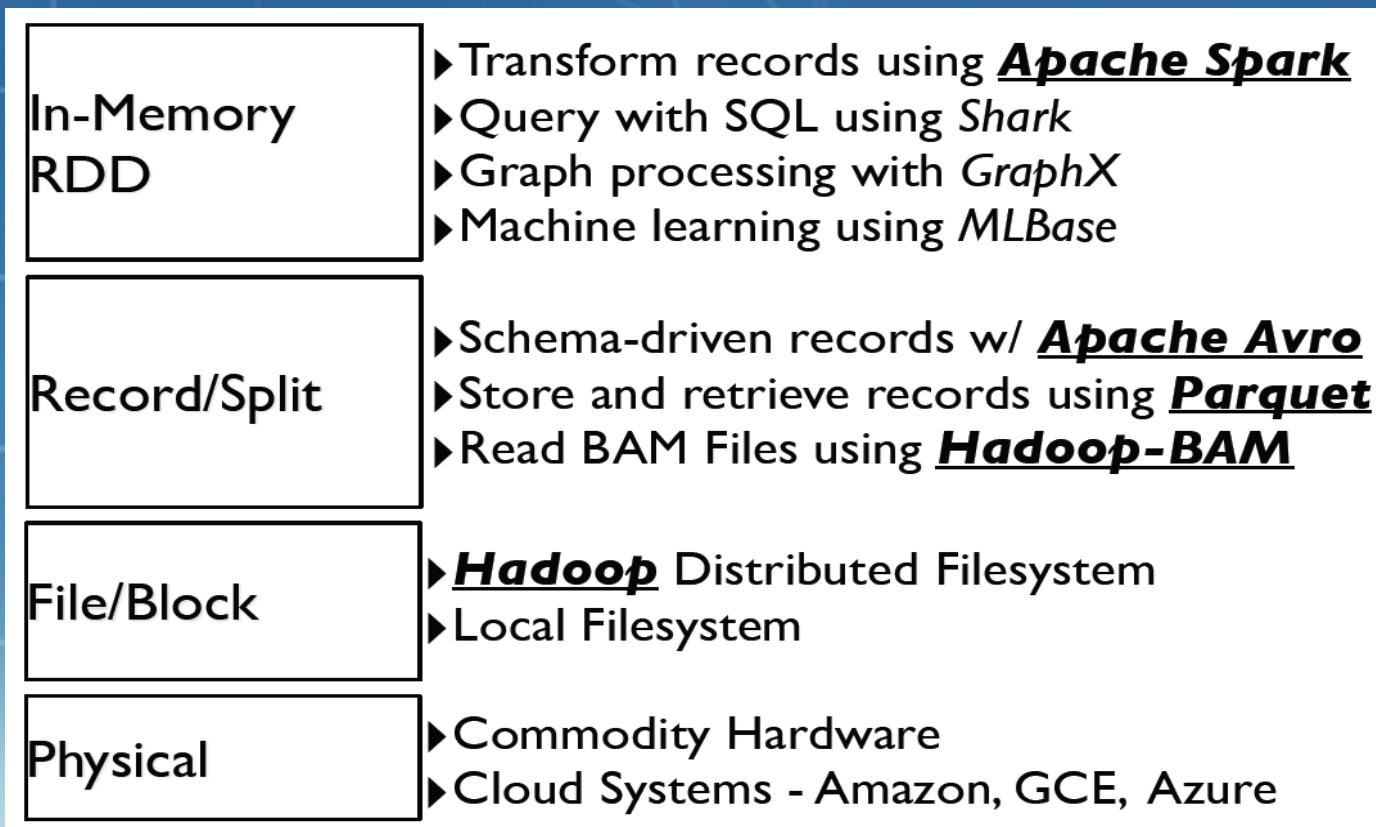
- Java library for the **manipulation** of files using the Hadoop MapReduce framework with the Picard SAM JDK
- BAM, SAM, FASTQ, FASTA, QSEQ, BCF, and VCF
- ADAM, SeqPig using
- Develop processing pipeline that enables efficient, scalable use of cluster/cloud
- Provide **data format** that has efficient parallel/distributed access across platforms

<http://sourceforge.net/projects/hadoop-bam/>



ADAM

A genomics processing engine and specialized file format built using Apache Avro, Apache Spark and **Parquet**.



ADAM Parquet

- Apache Parquet is a columnar storage format available to any project in the Hadoop ecosystem, regardless of the choice of data processing framework, data model or programming language.
- OSS Created by Twitter and Cloudera, based on Google Dremel, just entered Apache Incubator
- Columnar File Format:
 - Limits I/O to only data that is needed
 - Compresses very well - ADAM files are 5-25% smaller than BAM files without loss of data
 - Fast scans - load only columns you need, e.g. scan a read flag on a whole genome, highcoverage file in less than a minute



ADAM

Avocado

- A distributed pipeline for calling variants, and is built on top of Apache Spark and the ADAM API

ADAM:
Core API +
CLIs

bdg-formats:
Data schemas

avocado:
Distributed local
assembler

RNAAdam:
RNA analysis on
ADAM

bdg-services:
ADAM clusters

xASSEMBLEx:
GraphX-based de
novo assembler

Guacamole:
Distributed
somatic caller





SeqPig

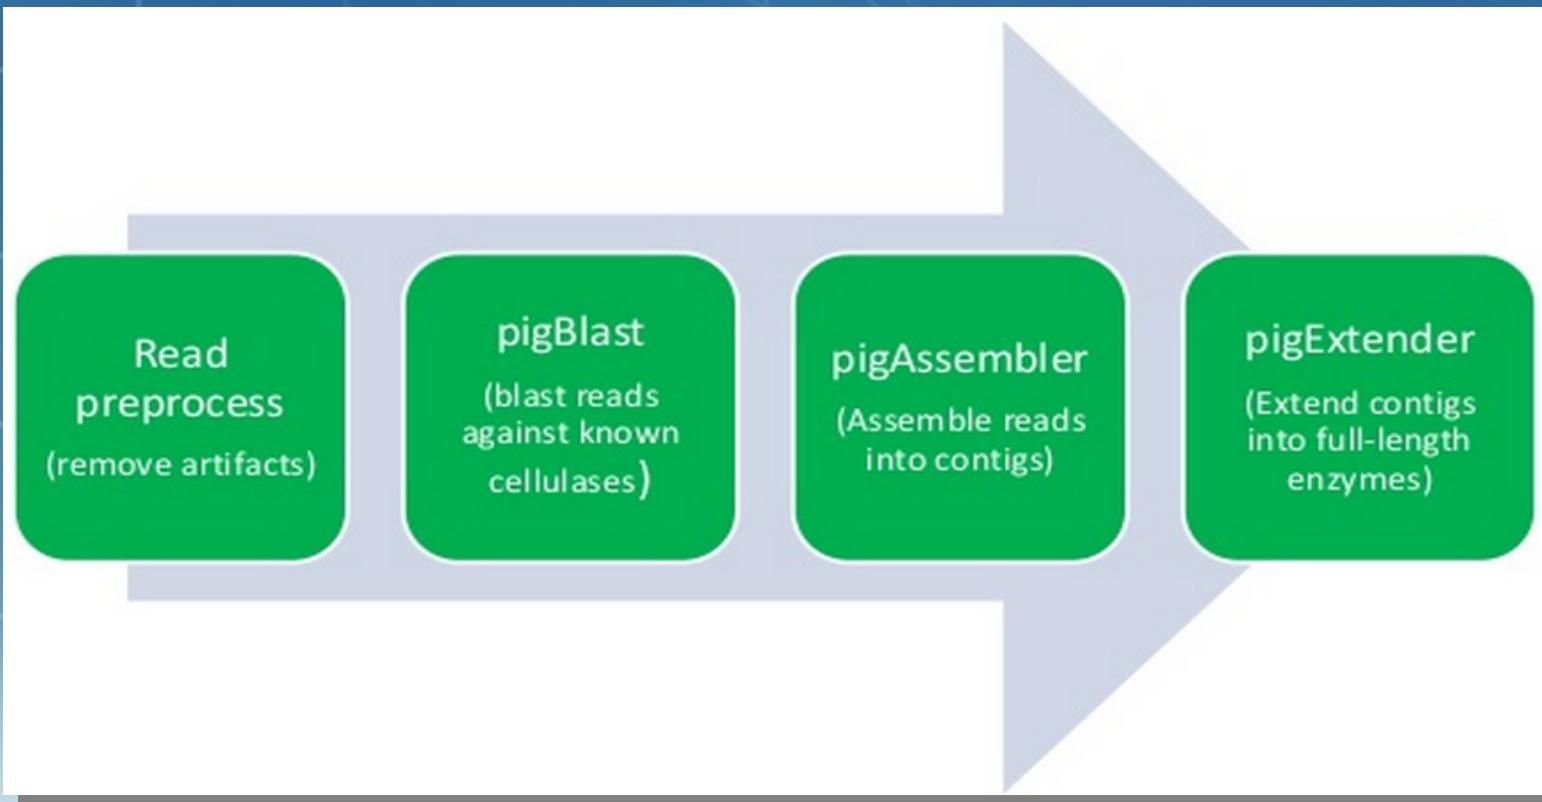
A library for Apache Pig for the distributed analysis
user-defined-functions (UDF's)
supports BAM/SAM, FastQ and Qseq input and output

<http://sourceforge.net/projects/seqpig/>



BioPig

- <https://github.com/JGI-Bioinformatics/biopig>
- <http://www.slideshare.net/ZhongWang3/biopig-for-large-sequencing-data>





Hadoop Seal

a suite of distributed applications for **aligning** short DNA reads, and **manipulating** and **analyzing** short read alignments.

Bcl2Qseq

Extract reads in qseq format from an Illumina run directory.

Demux

Separate sample data in the qseq file format produced by a multiplexed Illumina run.

PairReadsQSeq

convert files in the Illumina qseq file format to **prq format** to be processed by the alignment program, Seqal.

Seqal

Distributed short read mapping and duplicate removal tool.

ReadSort

Distributed **sorting** of read mappings.

RecabTable

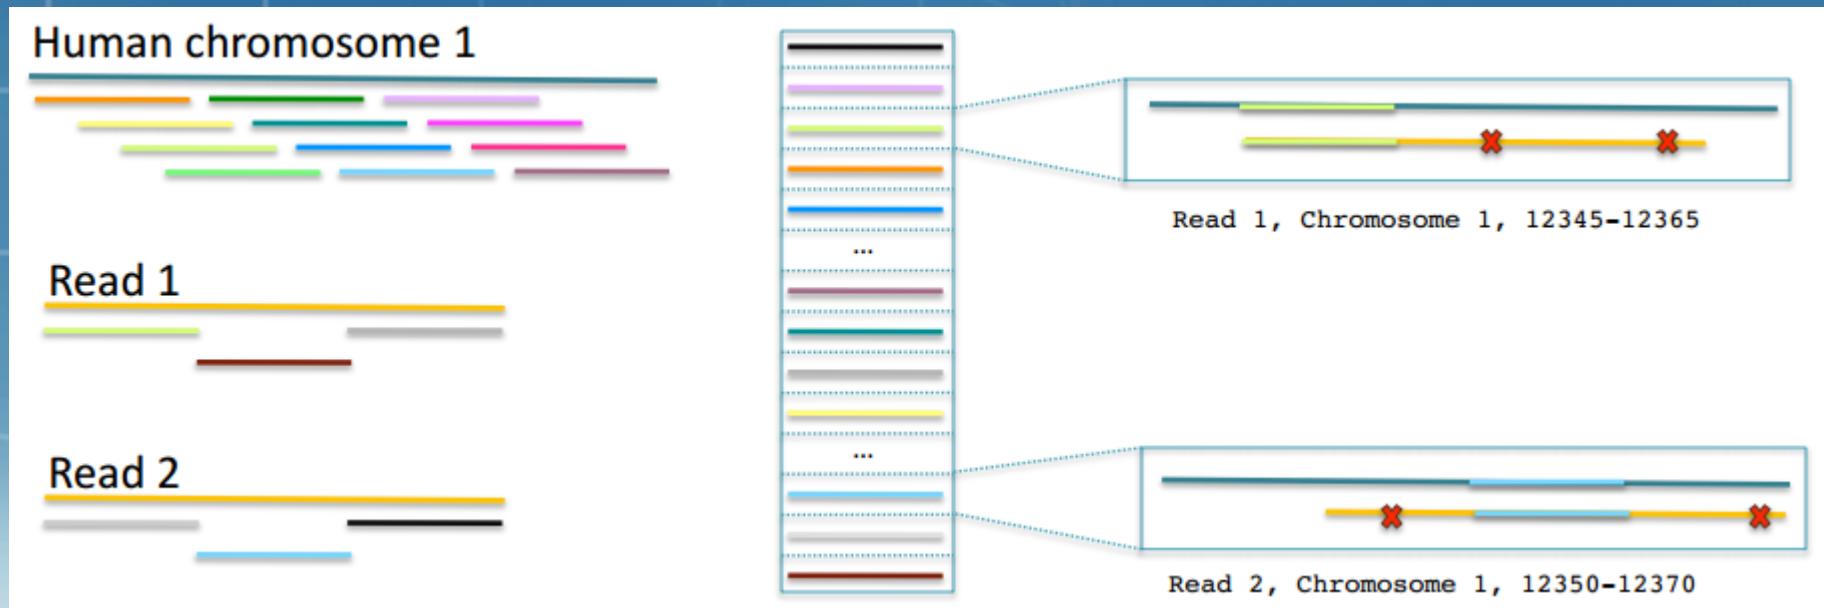
distributed **calculation** of covariates table to **estimate** empirical base qualities.



CloudBurst

- Highly Sensitive Short Read Mapping with MapReduce
- SNP discovery, genotyping, and personal genomics

July 8, 2010 - CloudBurst 1.1.0 released with support for Hadoop 0.20.X





Crossbow

- a scalable, portable, and automatic Cloud Computing tool for finding SNPs from short read data
- Genotyping from short reads using cloud computing
- combines Bowtie and SoapSNP
- Crossbow scripts require Perl 5.6 or later.
- Johns Hopkins Univ.



Outline

- ◆ Bigdata
- ◆ Hadoop
- ◆ Spark
- ◆ 3rd party
- ◆ **Machine Learning**
- ◆ Streaming
- ◆ Future





Apache Mahout

- Scalable machine learning library.
- Goodbye MapReduce !
- Mahout's Spark Shell Binding
- Algorithms
 - User and Item based recommenders
 - Matrix factorization based recommenders
 - K-Means, Fuzzy K-Means clustering
 - Latent Dirichlet Allocation
 - Singular Value Decomposition
 - Logistic regression classifier
 - (Complementary) Naive Bayes classifier
 - Random forest classifier
 - High performance java collections



H2O

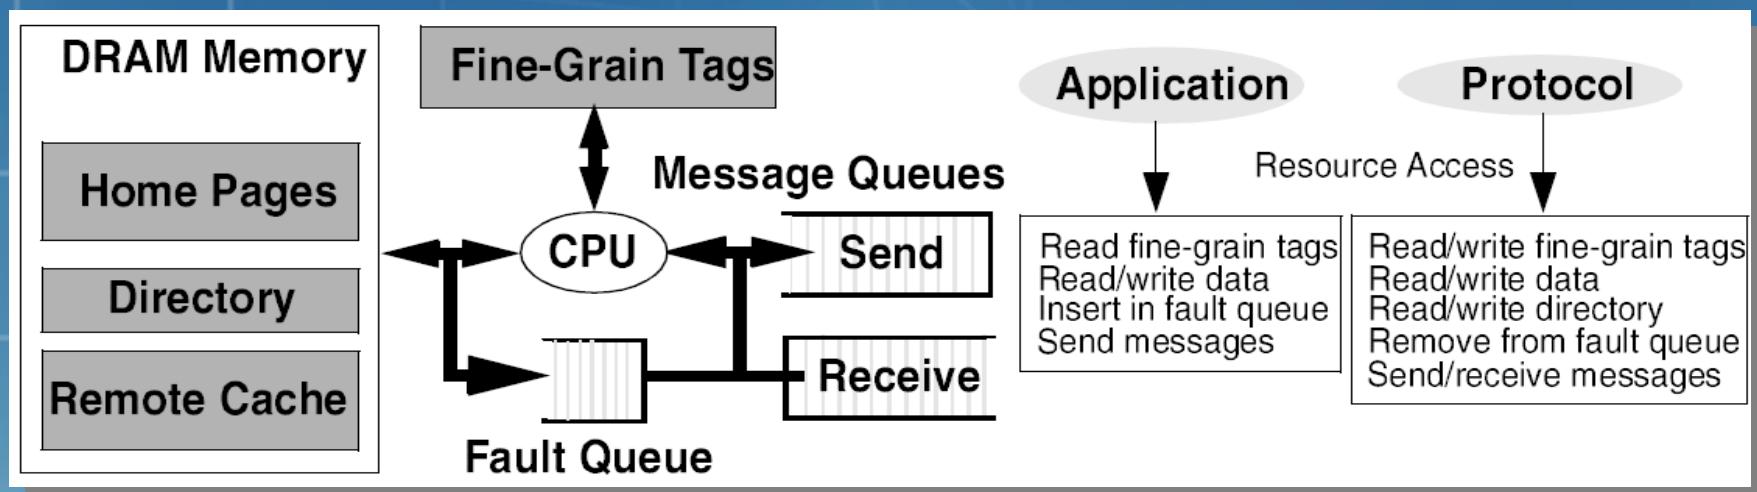
- The world's fastest in-memory platform for machine learning and predictive analytics on big data
- Fine-Grain Distributed Processing
- Support Deep Learning Neural Network using R
- Support for Java, Scala, and Python,R.
The solution's interface is driven by JSON APIs



H2O

Fine-Grain Distributed Processing

combines the responsiveness of in-memory processing with the ability to run fast **serialization between nodes and clusters**





Vowpal Wabbit

- a fast out-of-core learning system.
- Two ways to have a fast learning algorithm
 - start with a slow algorithm and speed it up
 - build an intrinsically fast learning algorithm
- by Microsoft Research and (previously) Yahoo! Research.

https://github.com/JohnLangford/vowpal_wabbit/wiki



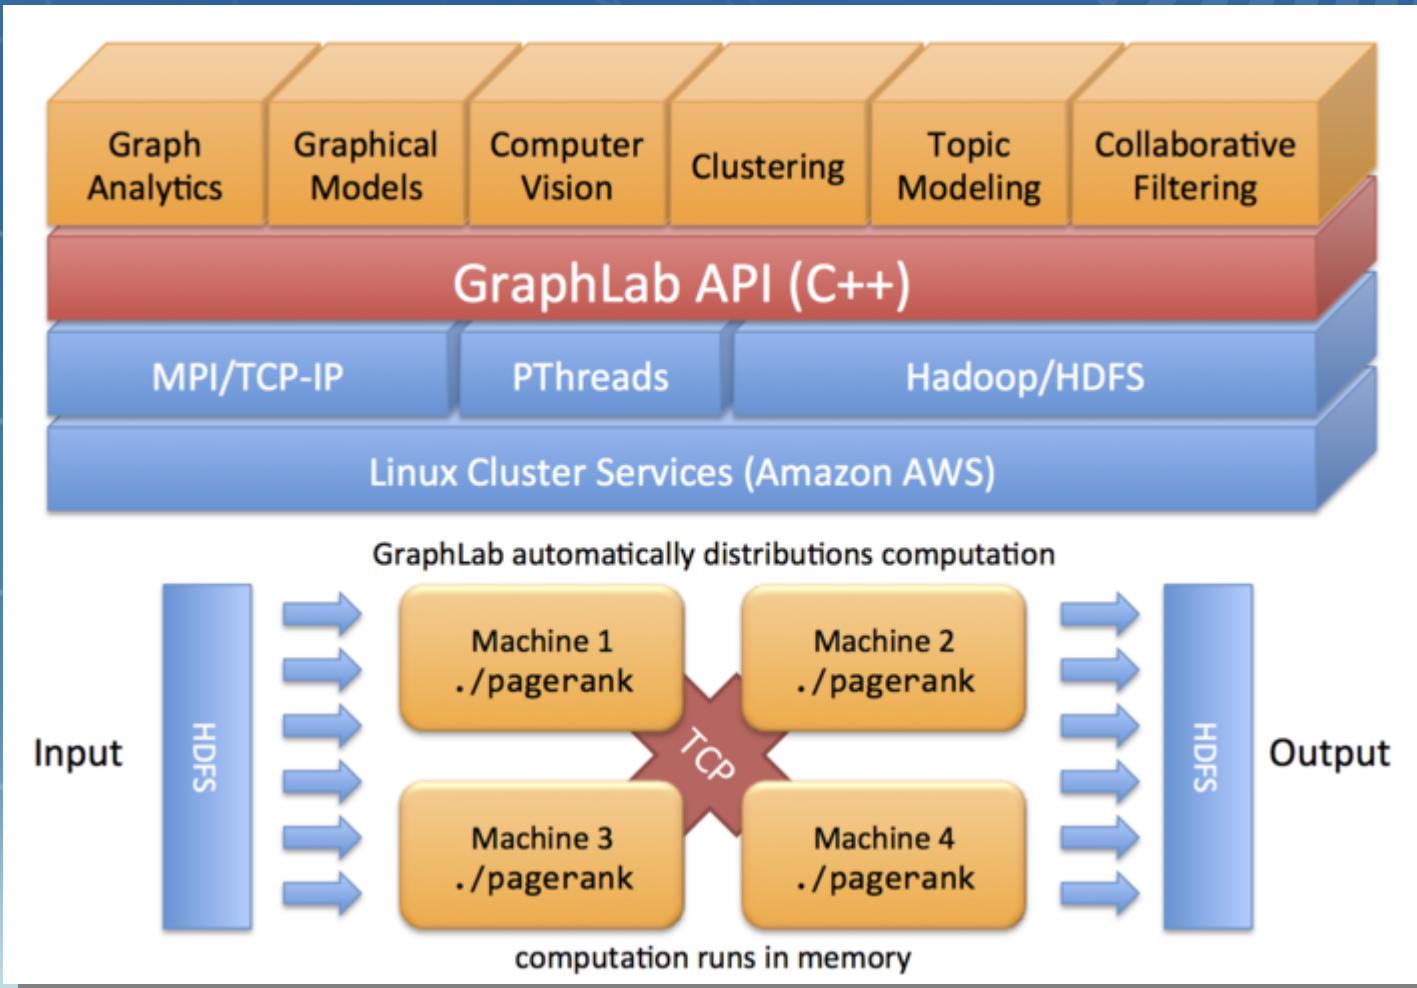
GraphLab

- Carnegie Mellon University in 2009 to develop a new parallel computation abstraction tailored to machine learning
- Toolkits
 - Topic Modeling
 - graph_algorithms
 - Graph Analytics
 - Clustering
 - Collaborative Filtering
 - Graphical Models
 - Factor Graph Toolkit
 - Linear iterative solver
 - Computer Vision



GraphLab

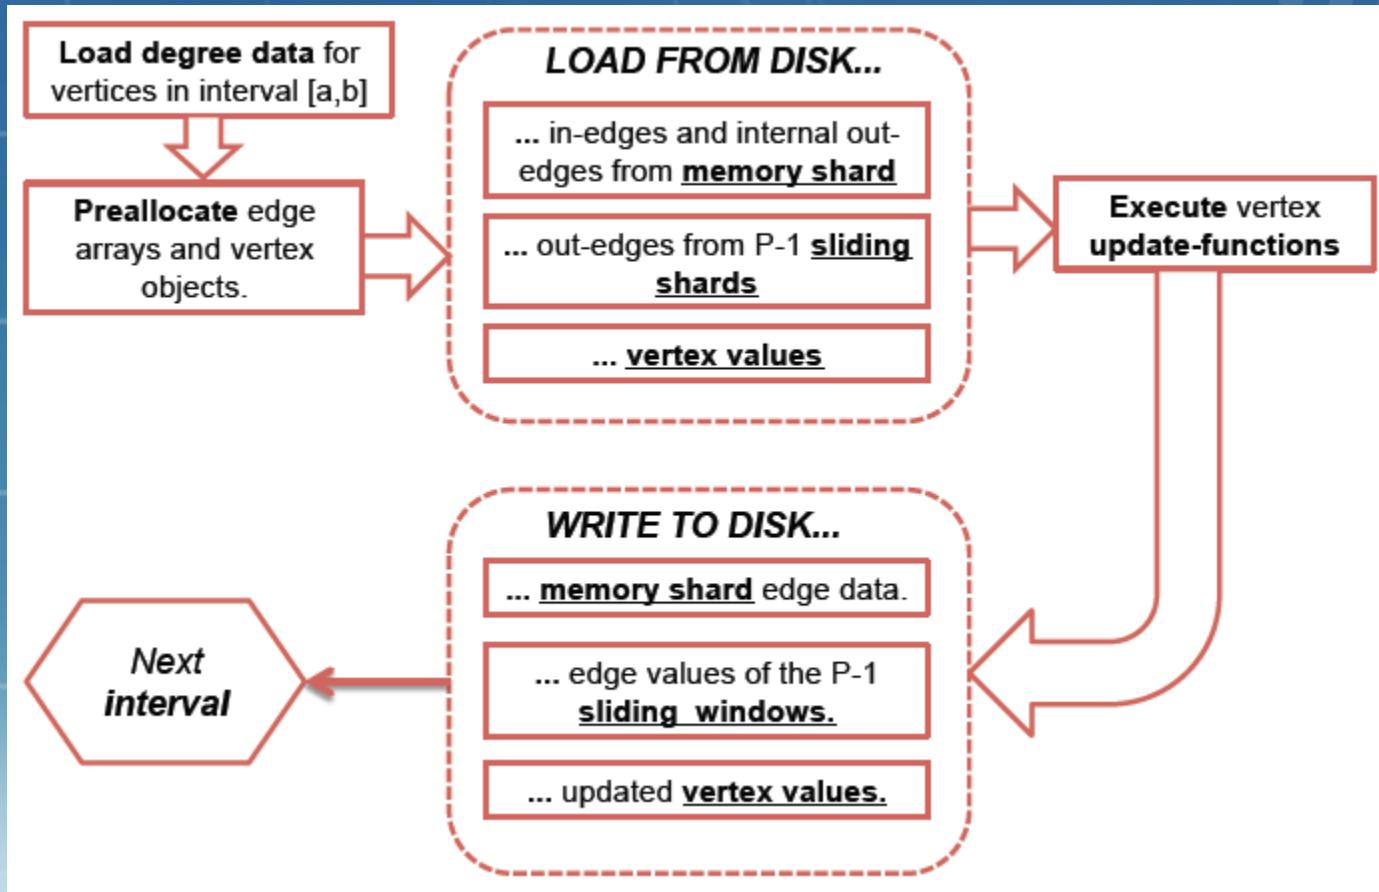
- Automatically distributions computation on HDFS





GraphLab GraphChi

Disk-based large-scale graph computation



Apache Giraph

An iterative graph processing system built for high scalability

Giraph originated as the open-source counterpart to Pregel

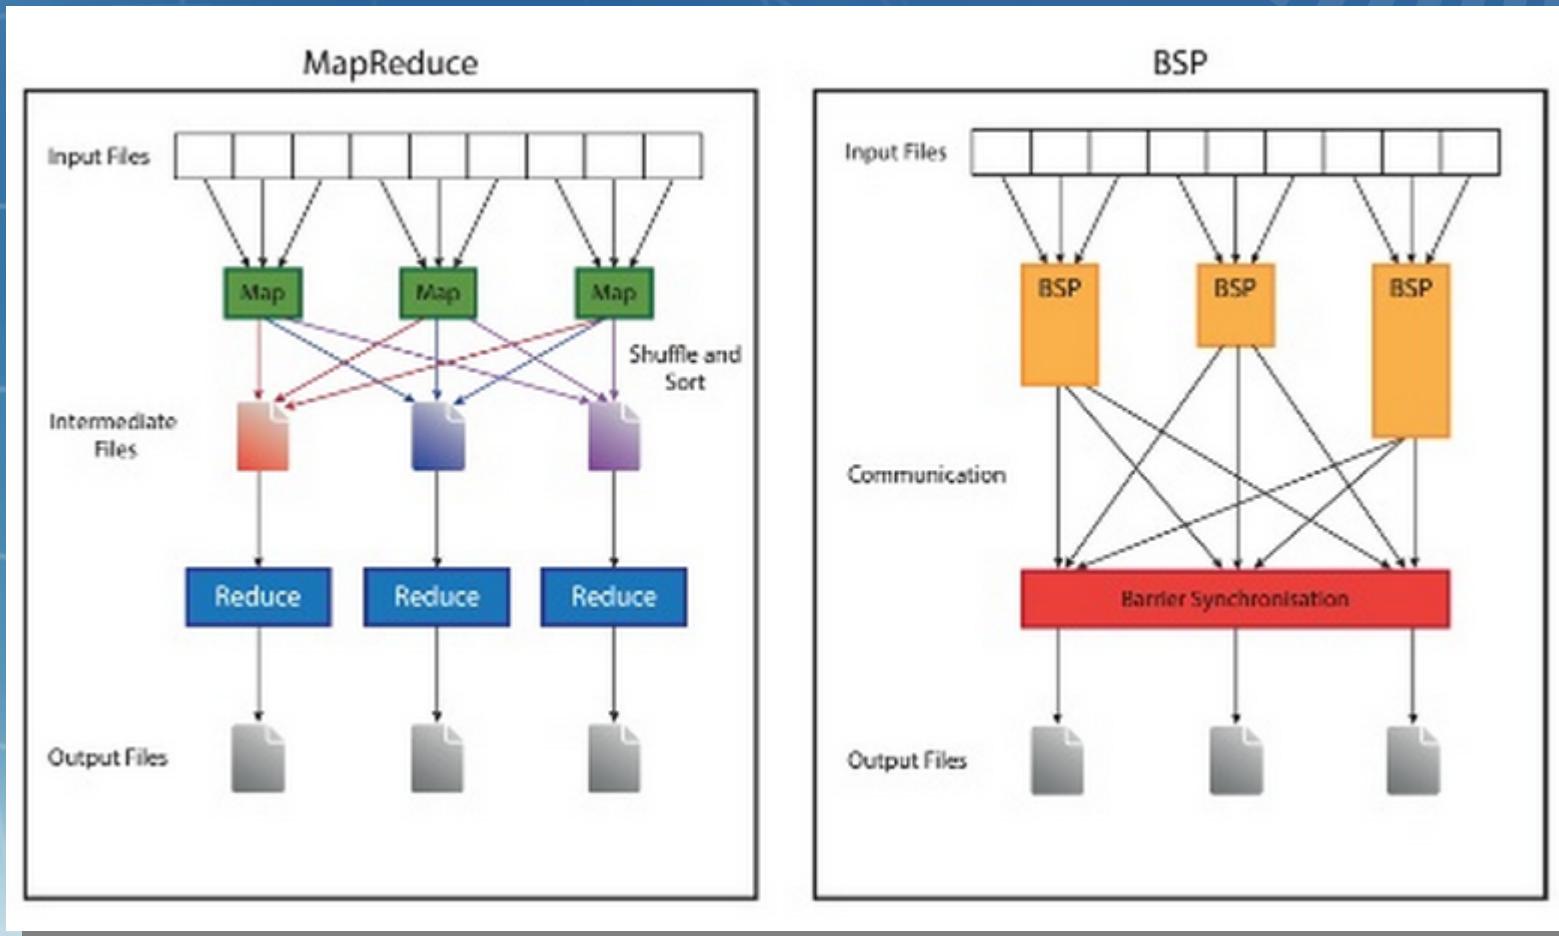
are inspired by the Bulk Synchronous Parallel model of distributed computation introduced by Leslie Valiant.





Apache Hama

A general BSP framework on top of Hadoop



Apache Hama

<http://www.slideshare.net/udanax/apache-hama-at-samsung-open-source-conference>

Apache Hama's **Advanced Analytics** Examples:

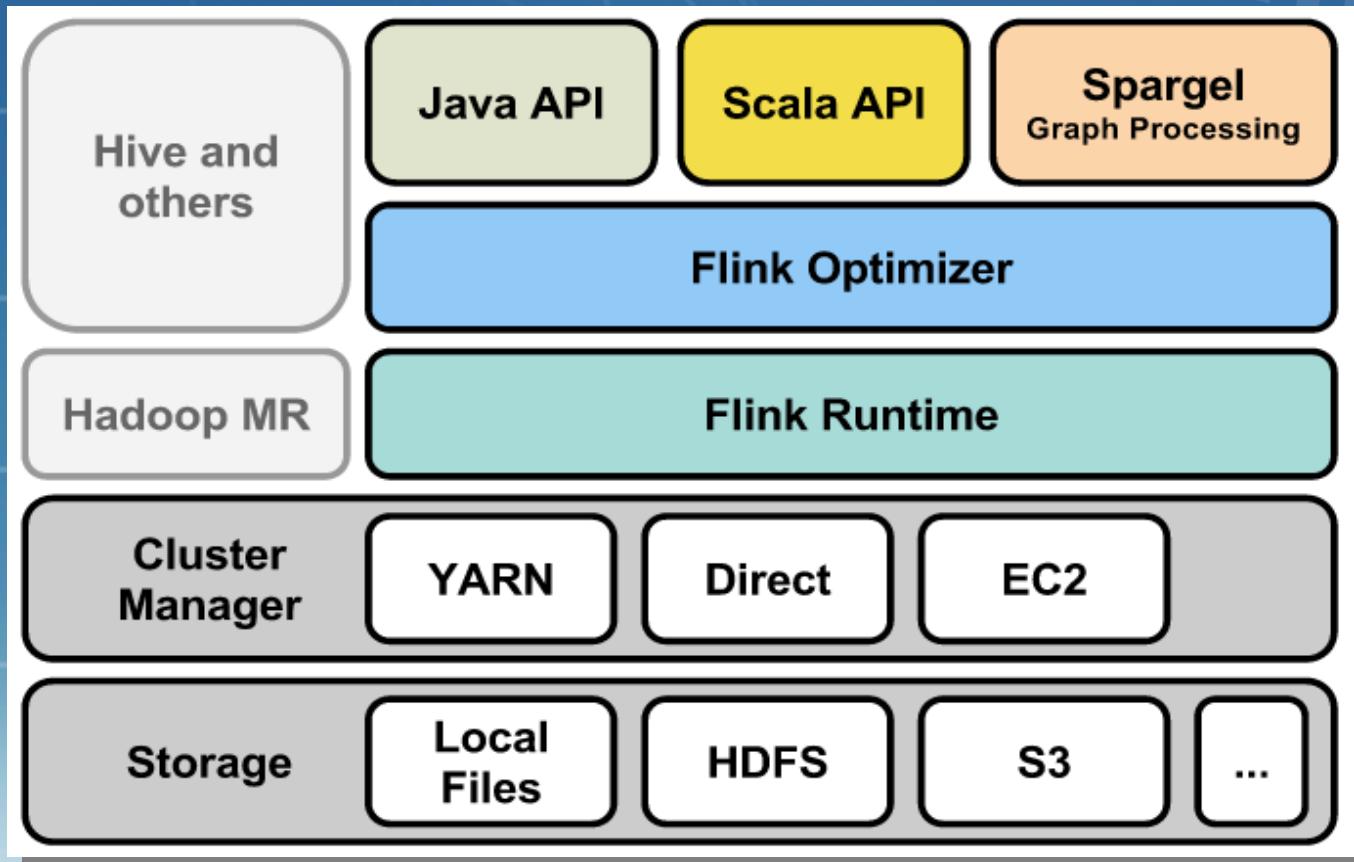
- Sparse Matrix-Vector Multiplication
- Semi-Clustering
- **K-means Clustering**
- Neural Networks
- **Gradient Descent**
- **PageRank**
- Single Source Shortest Path
- Bipartite Matching



Apache Flink(Stratosphere)

A platform for efficient, distributed, general-purpose data processing
Programming APIs for different languages (Java, Scala) and paradigms (record-oriented, graph-oriented).

Incubating..



Tajo

A big data warehouse system on Hadoop

Compatible

- ANSI/ISO SQL standard compliance
- Hive MetaStore access support
- JDBC driver support
- Various file formats support, such as CSV, RCFile, RowFile, SequenceFile and Parque



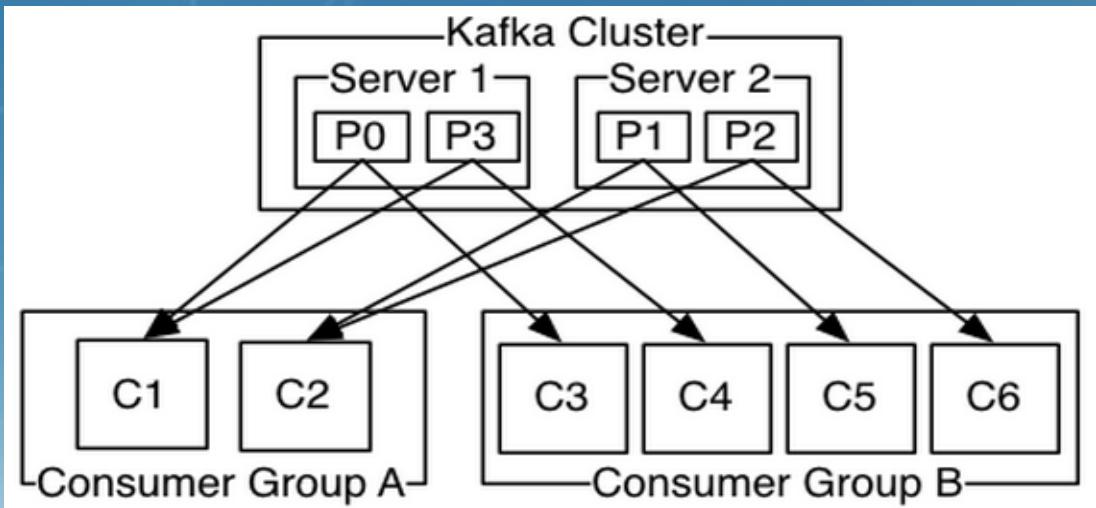
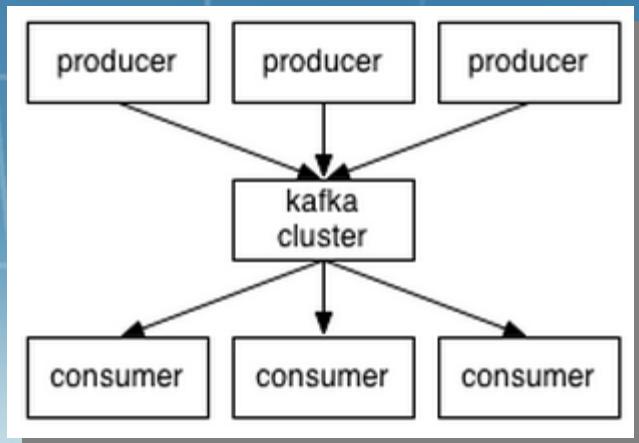
Outline

- ◆ Bigdata
- ◆ Hadoop
- ◆ Spark
- ◆ 3rd party
- ◆ Machine Learning
- ◆ **Streaming**
- ◆ Future



Kafka

- A high-throughput distributed messaging system.
- Producers
 - publish data to the topics of their choice.
- Consumers
 - Messaging traditionally has two models: queuing and publish-subscribe.



Kafka

Ecosystem:

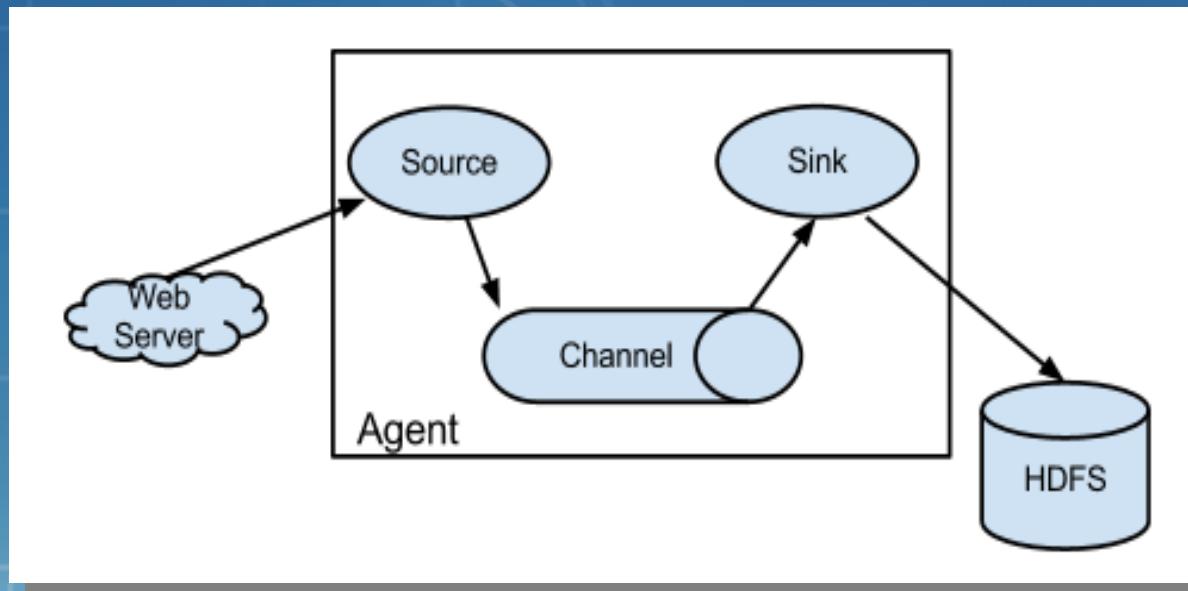
- Storm - A stream-processing framework.
- Samza - A YARN-based stream processing framework.
- Storm Spout - Consume messages from Kafka and emit as Storm tuples
- Kafka-Storm - Kafka 0.8, Storm 0.9, Avro integration
- SparkStreaming - Kafka receiver supports Kafka 0.8 and above
- Camus - LinkedIn's Kafka=>HDFS pipeline. This one is used for all data at LinkedIn, and works great.
- Kafka Hadoop Loader A different take on Hadoop loading functionality from what is included in the main distribution.
- Flume - Contains Kafka Source (consumer) and Sink (producer)



Flume

A distributed, reliable, and available service for efficiently collecting, aggregating, and moving large amounts of log data.

Avro
Thrift
Syslog
Netcat

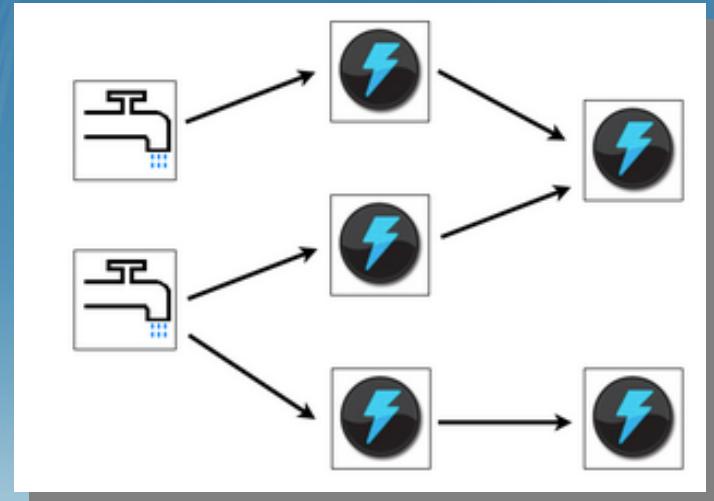
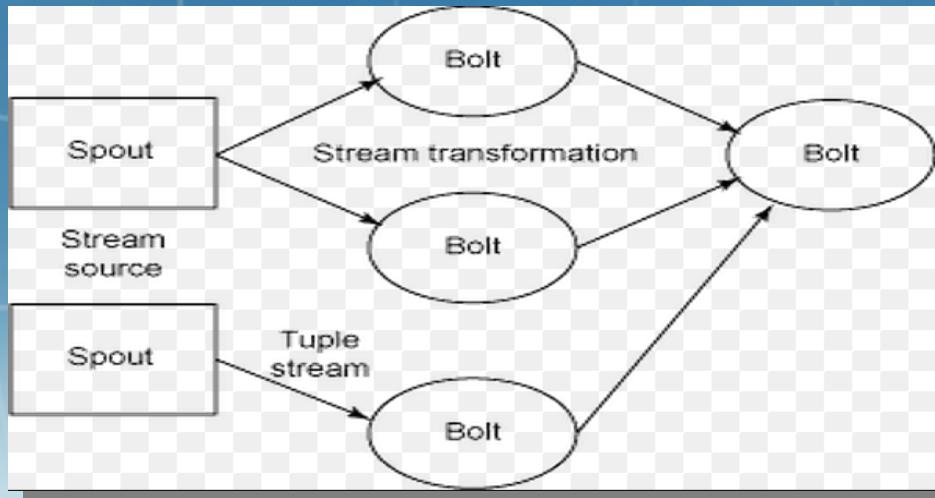


Storm

a free and open source distributed realtime computation system

Topologies

To do realtime computation on Storm, you create what are called “topologies”. A topology is a graph of computation.



MQTT(MQ Telemetry Transport)

A machine-to-machine (M2M)/"Internet of Things" connectivity protocol.

- publish/subscribe messaging transport
- used in sensors communicating to a broker via satellite link, over occasional dial-up connections with healthcare providers

Three qualities of service for message delivery:

- "At most once", where messages are delivered according to the best efforts of the underlying TCP/IP network. Message loss or duplication can occur. This level could be used, for example, with ambient sensor data where it does not matter if an individual reading is lost as the next one will be published soon after.
- "At least once", where messages are assured to arrive but duplicates may occur.
- "Exactly once", where message are assured to arrive exactly once. This level could be used, for example, with billing systems where duplicate or lost messages could lead to incorrect charges being applied.

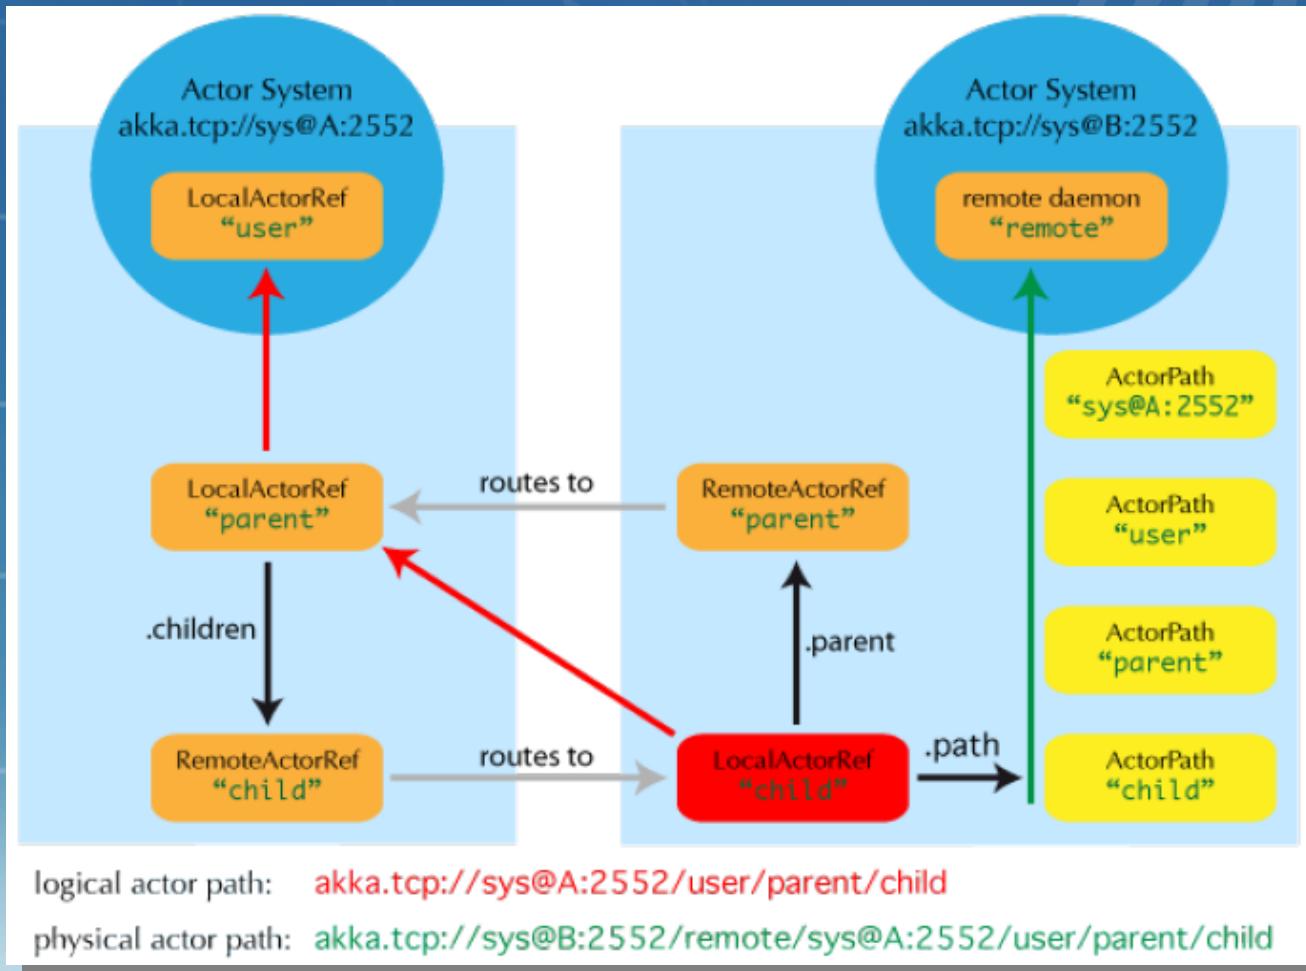


Akka

- Scalable real-time transaction processing
- an unified runtime and programming model
 - Scale up (Concurrency)
 - Scale out (Remoting)
 - Fault tolerance
- Actor
 - Simple and high-level abstractions for concurrency and parallelism.
 - Asynchronous, non-blocking and highly performant event-driven programming model.
 - Very lightweight event-driven processes (several million actors per GB of heap memory).

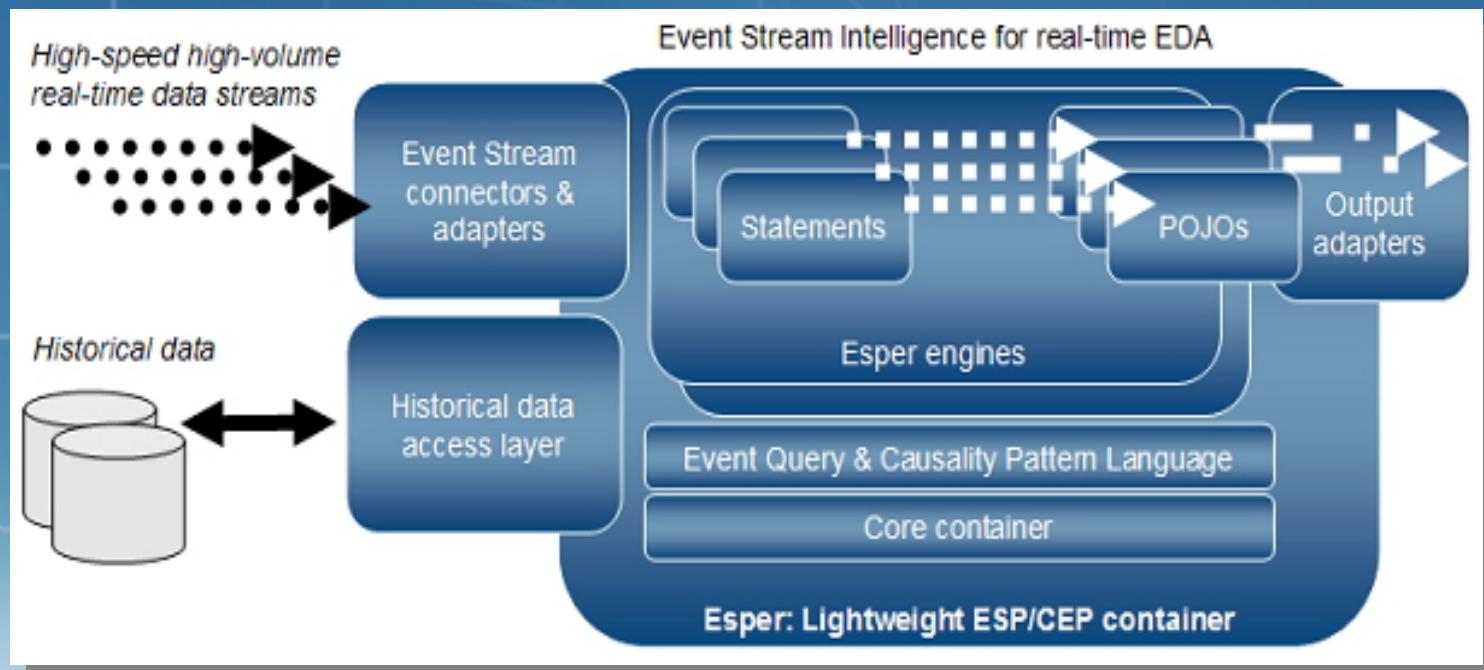


Akka Remote Deployment



Esper

- a component for complex event processing (CEP) and event series analysis
- Event Processing Language (EPL)
 - highly scalable, memory-efficient, in-memory computing, SQL-standard, minimal latency, real-time streaming-capable
- offers a Domain Specific Language (DSL) for processing events.



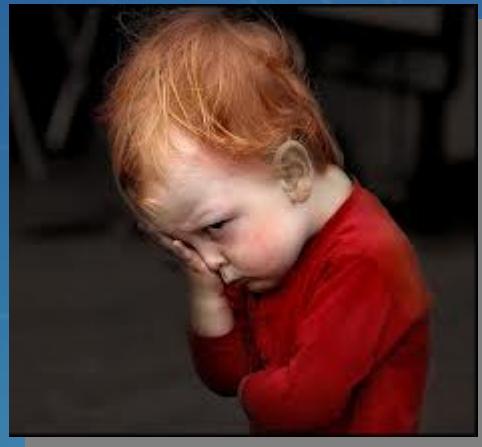
Outline

- ◆ Bigdata
- ◆ Hadoop
- ◆ Spark
- ◆ 3rd party
- ◆ Machine Learning
- ◆ Streaming
- ◆ **Future**



Bigdata and Machine Learning

Very Difficult !



Solution

- For End User
- For ML Developer
- Interactive Advanced Machine Learning Development Platform
- DSL style interactive ML modeling(like matlab)
- Domain specific libraries and tool



Conclusion

- Bigdata is very important
- ML needs Computation power (Graphical model)
- Real time streaming and computation
- High throughput Streaming
- Data size will improve the ML
- Biobigdata needs scalable ML
- High throughput sequence analysis
- De novo Genome Assembly is base step

BIcube

Manipul

Mashup

Preprocess

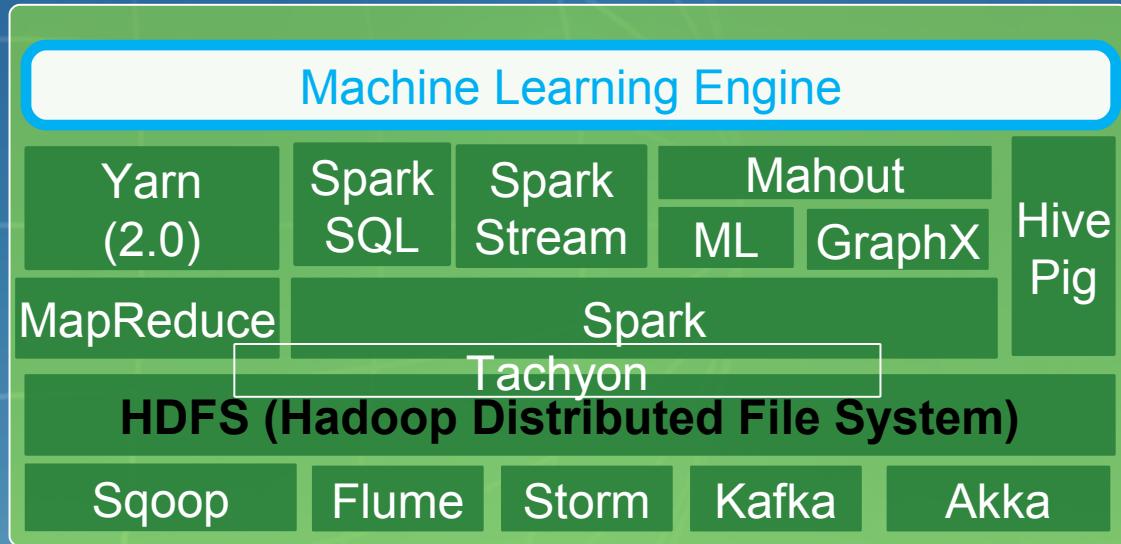
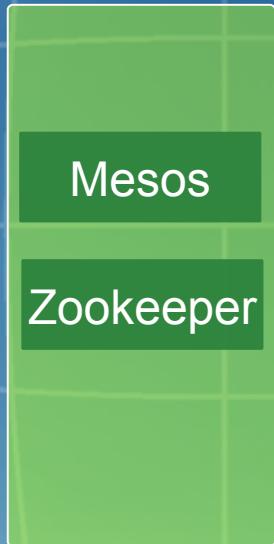
Classify

Cluster

Associate

Bioinfo

Visual



RDB Mashup

SNS Crawler

Log Crawler

Authentication

CubeManager



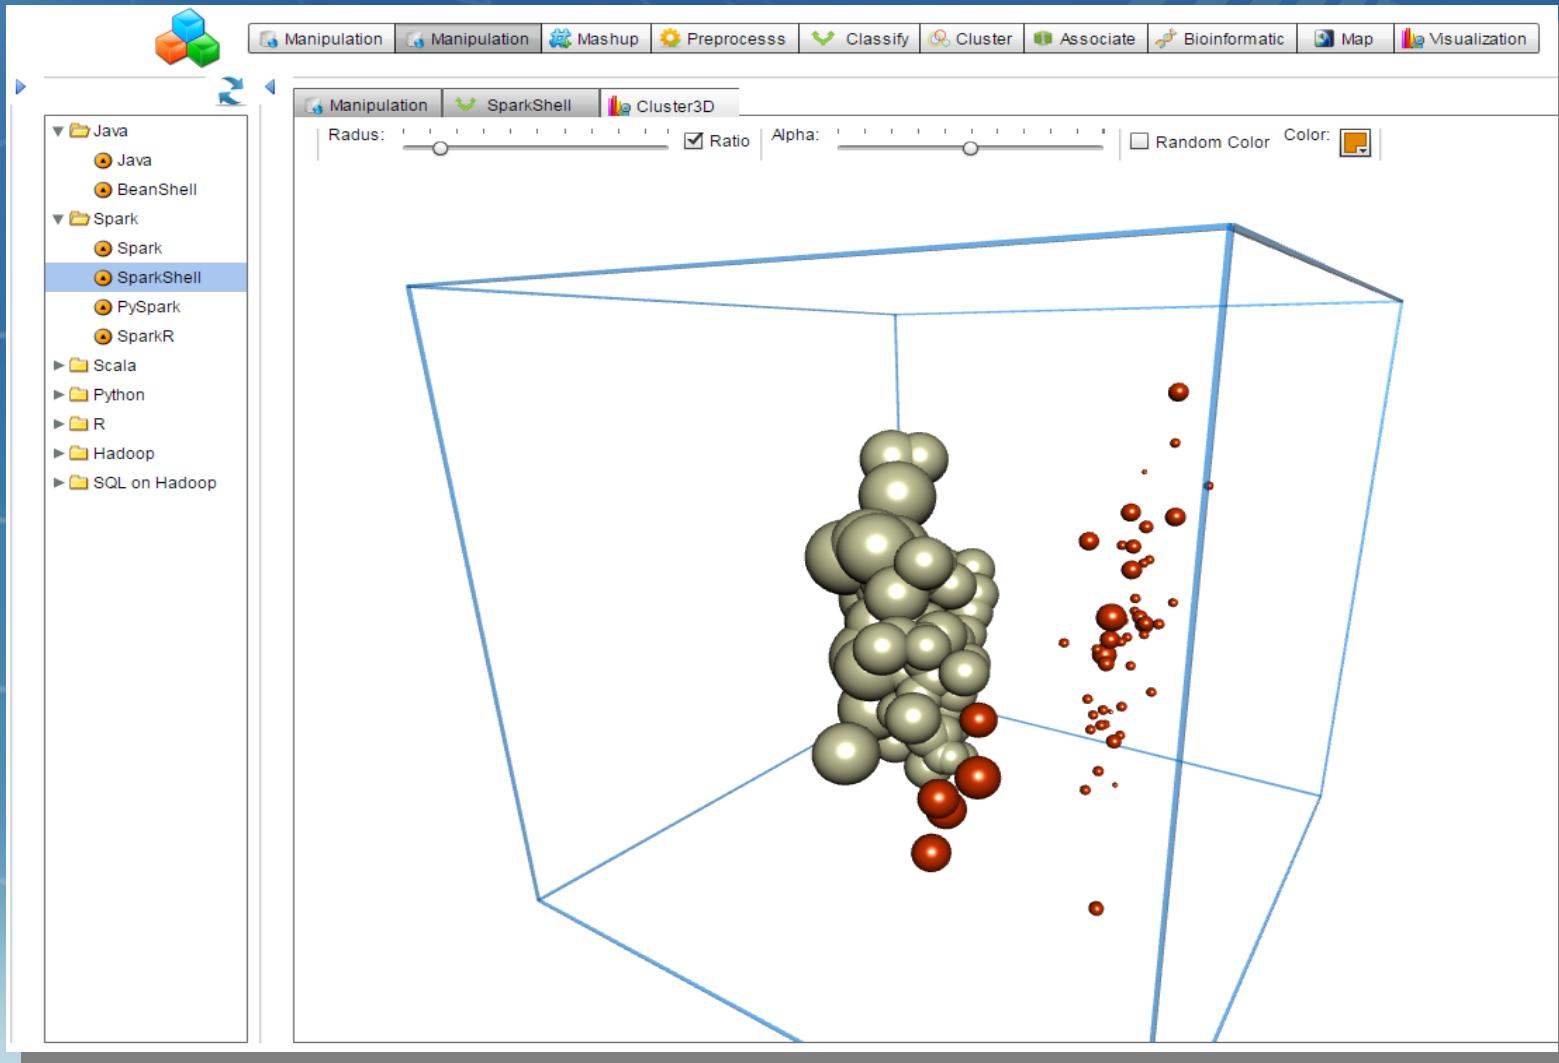
Spark Interactive Shell

The screenshot shows the Bio Bigdata application interface with the "Manipulation" tab selected. On the left, there is a navigation tree with categories like Java, Spark, Scala, Python, R, Hadoop, and SQL on Hadoop. Under the "Spark" category, "SparkShell" is selected. The main window displays the output of a Spark session. It starts with a welcome message and version information, followed by log entries from the sparkDriver and sparkWorker processes. The log includes details about the driver starting remoting, creating local directories, and successfully starting services.

```
28 Welcome to
29
30   _/\_   _/\_   _/\_
31  \ V_ V_ 'V_ V_ '
32  /_/_/_\_,/_/_/_\_
33  /_/
34
35 Using Scala version 2.10.4 (Java HotSpot(TM) 64-Bit Server VM, Java 1.8.0_05)
36 Type in expressions to have them evaluated.
37 Type :help for more information.
38 14/09/30 14:58:19 WARN Utils: Your hostname, studyand2 resolves to a loopback address:
39 127.0.0.1; using 1.234.27.234 instead (on interface eth0)
40 14/09/30 14:58:19 WARN Utils: Set SPARK_LOCAL_IP if you need to bind to another address
41 14/09/30 14:58:19 INFO SecurityManager: Changing view acls to: spark,
42 14/09/30 14:58:19 INFO SecurityManager: Changing modify acls to: spark,
43 14/09/30 14:58:19 INFO SecurityManager: SecurityManager: authentication disabled; ui acls
44 disabled; users with view permissions: Set(spark, ); users with modify permissions: Set(spark, )
45 14/09/30 14:58:20 INFO Slf4jLogger: Slf4jLogger started
46 14/09/30 14:58:20 INFO Remoting: Starting remoting
47 14/09/30 14:58:21 INFO Remoting: Remoting started; listening on addresses :[akka.tcp://
48 sparkDriver@1.234.27.234:38090]
49 14/09/30 14:58:21 INFO Remoting: Remoting now listens on addresses: [akka.tcp://
50 sparkDriver@1.234.27.234:38090]
51 14/09/30 14:58:21 INFO Utils: Successfully started service 'sparkDriver' on port 38090.
52 14/09/30 14:58:21 INFO SparkEnv: Registering MapOutputTracker
53 14/09/30 14:58:21 INFO SparkEnv: Registering BlockManagerMaster
54 14/09/30 14:58:21 INFO DiskBlockManager: Created local directory at /tmp/spark-
55 local-20140930145821-180c
56 14/09/30 14:58:21 INFO Utils: Successfully started service 'Connection manager for block
manager' on port 50734.
```

command>

View GPU Cluster3D using Spark Interactive Shell



Thanks !
Question ?

and Demo...



```
import com.chimpler.example.kmeans._  
import com.bicube.akka.visualization.VisualDeploy  
  
val args2 = Array(  
    "/user/root/spark/data/twitter.location"  
    , "/user/root/spark/data/kmean.image3.png"  
)  
  
KMeansAppHadoop.run(args2, sc)  
  
  
import com.bicube.akka.visualization.VisualDeploy  
val v=new VisualDeploy()  
v.deploy("{ " +  
"\\"name\":\\"ImageViewer\\\" " +  
", \\"module\":\\"image\\\" " +  
", \\"path\":\\"spark/data/kmean.image3.png\\\" " +  
"}");
```



Demo

```
import com.bicube.akka.visualization._  
val v=new VisualDeploy()
```

```
v.deploy("{ " +  
" \"name\": \"PDBViewer\" " +  
", \"module\": \"pdb3d\" " +  
", \"path\": \"hdfs: /user/root/spark/data/1CRN.pdb.gz\" " +  
"}");
```

```
import com.bicube.akka.visualization._  
val v=new VisualDeploy()  
v.deploy("{ " +  
" \"name\": \"Cluster3D\" " +  
", \"module\": \"cluster3d\" " +  
", \"path\": \"hdfs: /user/root/cluster/kmeans-iris/output-iris\" " +  
"}");
```

